


Parallel Sparse Matrix Solver for Circuit Simulations using FPGA

Yogesh Mahajan
Supervisor: Prof. Sachin Patkar


Electrical Engineering, IIT Bombay
October 2018

Motivation

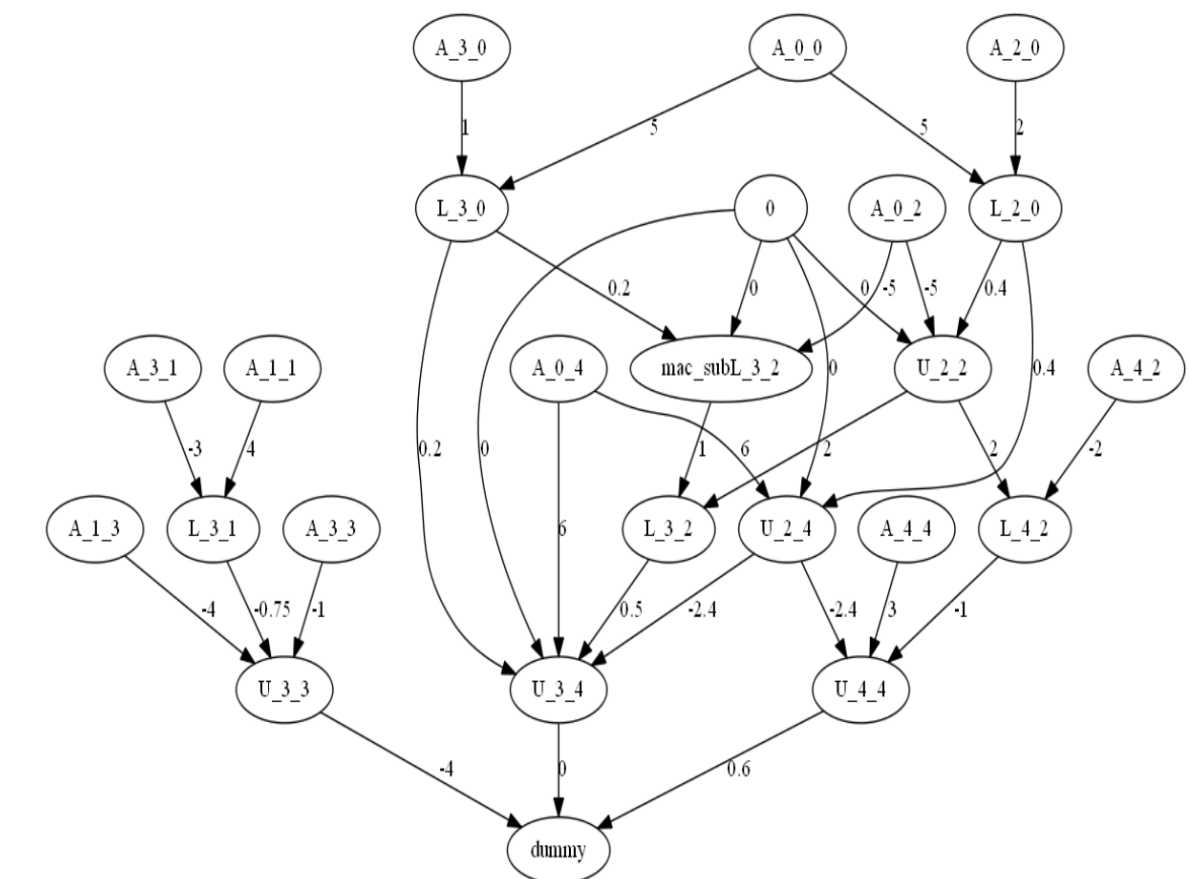
Why and How to Accelerate Circuit Simulations?



Simulation requires several thousand repeated solution of the network matrix



Structure of matrix remains the same in each iteration



Data manipulation can be modelled using a task graph to extract parallelism

Parallelizing the Sparse LU Decomposition

Pre-processing

1

- Convert matrix to sparse storage format
- Pre-order the matrix for fill in reduction
- Generate verification data

Symbolic Analysis

2

- Find the non-zero location in factors
- Generate the computation flow graph
- Assign priority to each operation

Scheduling

3

- Generate memory map
- Generate priority list based schedule for the computation flow graph

Numeric Computation

4

- Copy matrix data to FPGA
- Execute using the predefined schedule
- Write results back to main memory

Sparse Matrix Storage Formats

$$\begin{bmatrix} 5 & 0 & -5 & 0 & 6 \\ 0 & 4 & 0 & -4 & 0 \\ 2 & 0 & 0 & 0 & 0 \\ 1 & -3 & 0 & -1 & 0 \\ 0 & 0 & -2 & 0 & 3 \end{bmatrix}$$

Sparse Matrix

$$\begin{bmatrix} 5 & -5 & 6 \\ 4 & -4 & 0 \\ 2 & 0 & 0 \\ 1 & -3 & -1 \\ -2 & 3 & 0 \end{bmatrix} \begin{bmatrix} 0 & 2 & 4 \\ 1 & 3 & -1 \\ 0 & -1 & -1 \\ 0 & 1 & 3 \\ 2 & 4 & -1 \end{bmatrix}$$

ELLPACK

Values	5	-5	6	4	-4	2	-	-3	-1	-2	3
Column Indices	0	2	4	1	3	0	0	1	3	2	4
Row Indices	0	0	0	1	1	2	3	3	3	4	4

Triplet Format

Values	5	2	1	4	-3	-5	-2	-4	-1	6	3
Row Indices	0	2	3	1	3	0	4	1	3	0	4
Column Pointers	0			3		5		7		9	

Compressed Column Format

Values	5	-5	6	4	-4	2	-	-3	-1	-2	3
Column Indices	0	2	4	1	3	0	0	1	3	2	4
Row Pointers	0			3		5		6		9	

Compressed Row Format


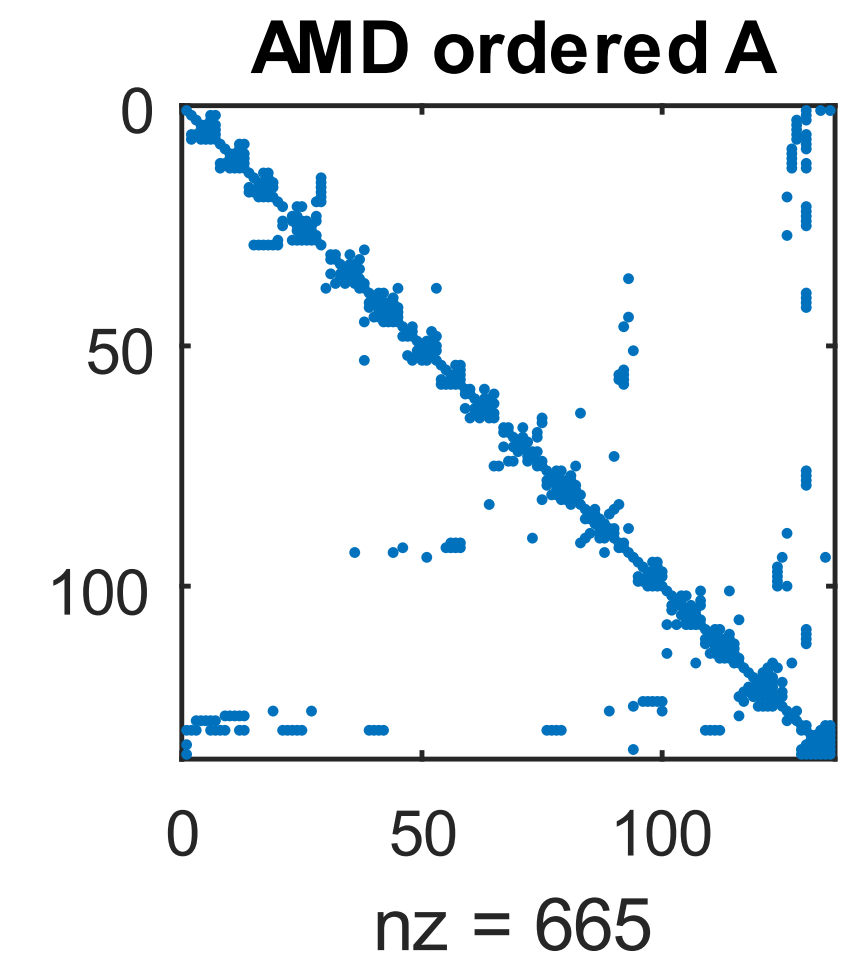
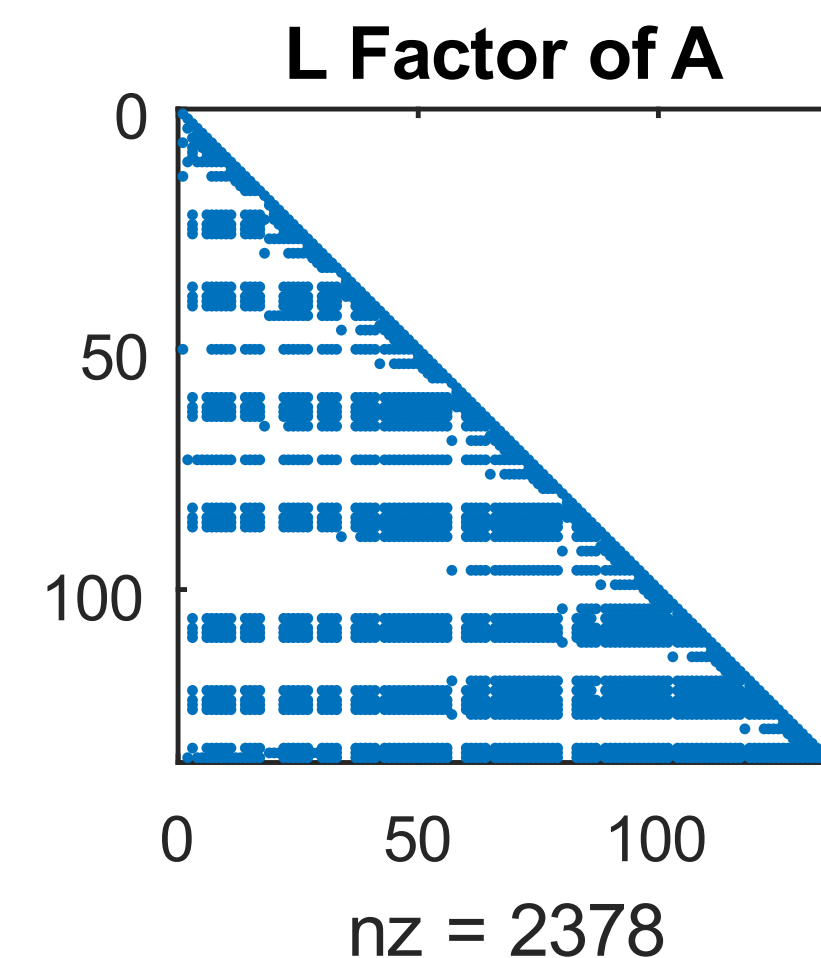

Approximate Minimum Degree Ordering

Find a permutation matrix $P = amd(A)$ such that the *Cholesky factor* of $P^T A P$ has less number of fill-ins than the factors of A , where A is a $n \times n$ symmetric matrix

For asymmetric matrix $P = amd(A + A^T)$

The linear problem $Ax = b$ can be solved by solving the reordered system


$$(P^T A P)(P^T x) = P^T b$$




Symbolic Analysis

$$U_{(i,j)} = A_{(i,j)} - \sum_{k=1}^{i-1} L_{(i,k)} U_{(k,j)} \qquad L_{(i,j)} = \frac{A_{(i,j)} - \sum_{k=1}^{j-1} L_{(i,k)} U_{(k,j)}}{U_{(j,j)}}$$


Crout's



Left-Looking



Right-Looking



Gilbert-Peierls' Algorithm


Algorithm 1 Gilbert-Peierls Algorithm: A Column-Oriented LU Factorization

Precondition: A , a $n \times n$ asymmetric matrix

```

1  $L := I$ 
2 for  $j := 1$  to  $n$  do
3   Solve  $L_j u_j = a_j$  for  $u_j$ 
4    $b'_j := a'_j - L'_j u_j$ 
5   Do Partial Pivoting on  $b'_j$ 
6    $u_{jj} := b_{jj}$ 
7    $l'_j := b'_j / u_{jj}$ 

```



J. Gilbert and T. Peierls, "Sparse partial pivoting in time proportional to arithmetic operations," SIAM Journal on Scientific and Statistical Computing, vol. 9, no. 5, pp. 862-874, 1988.

Solving the Lower Triangular System


Algorithm 2 Gilbert-Peierls Algorithm: Solving Triangular System $L_j x = b$

Precondition: L_j is a lower triangular matrix, x, b are sparse column vectors

```

1  $x := b$ 
2 for each  $j \in \chi$  do
3   for each  $i > j$  for which  $l_{ij} \neq 0$  do
4      $x_i := x_i - l_{ij}x_j$ 

```



$$(b_i \neq 0) \implies (x_i \neq 0)$$

$$(x_j \neq 0) \wedge (l_{ij} \neq 0) \implies (x_i \neq 0)$$

Finding Non-Zero Locations

Column dependence in the factorization process can be modelled using a DGA of columns.


Computational dependence for j^{th} column is given by *Reach* of that column in column dependence DAG:

$$Reach(j) = \bigcup_{i \in \{i | x_{ij} \neq 0\}} Reach(i)$$

Hence the fill-ins are given by:




$$fillIn(j) = Reach(j) - \{i | x_{ij} \neq 0\}$$

1							●		
	2		●						
●		3			●				
			4		●			●	
	●		●	5			●		
				●	6				
●		●			●	7			
							8		
			●					9	
							●		10





Finding Non-Zero Locations(Example)

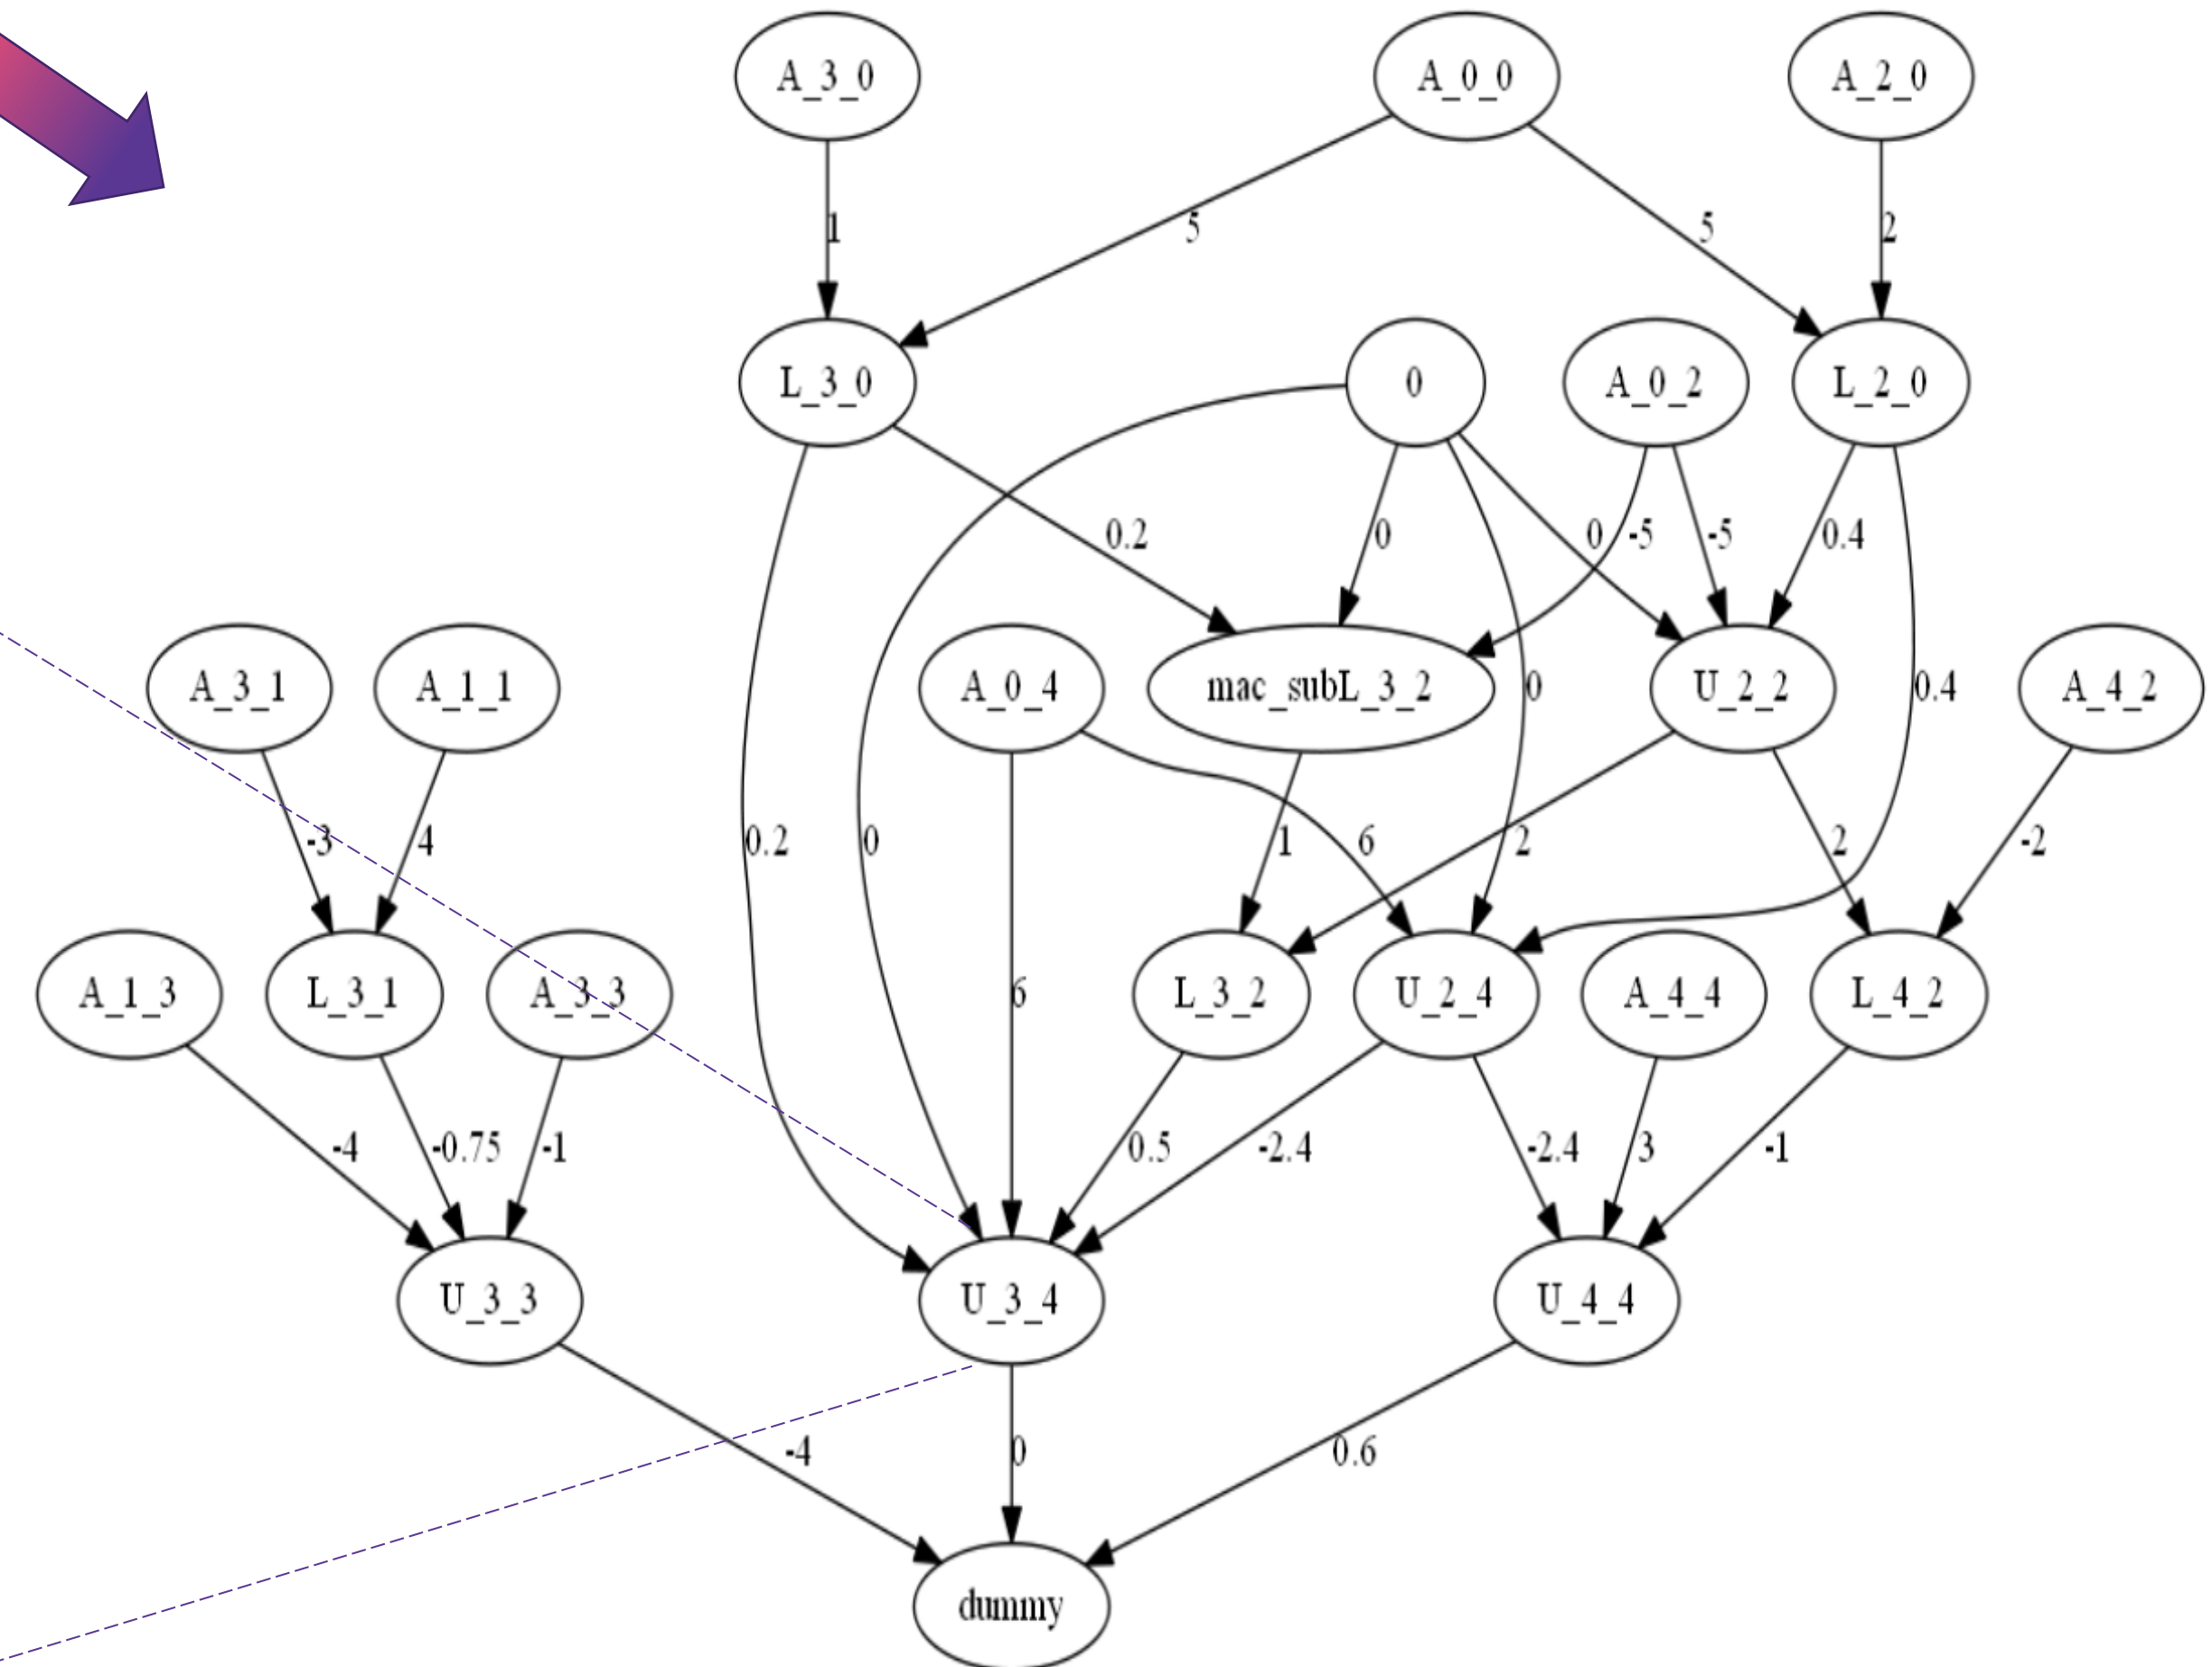
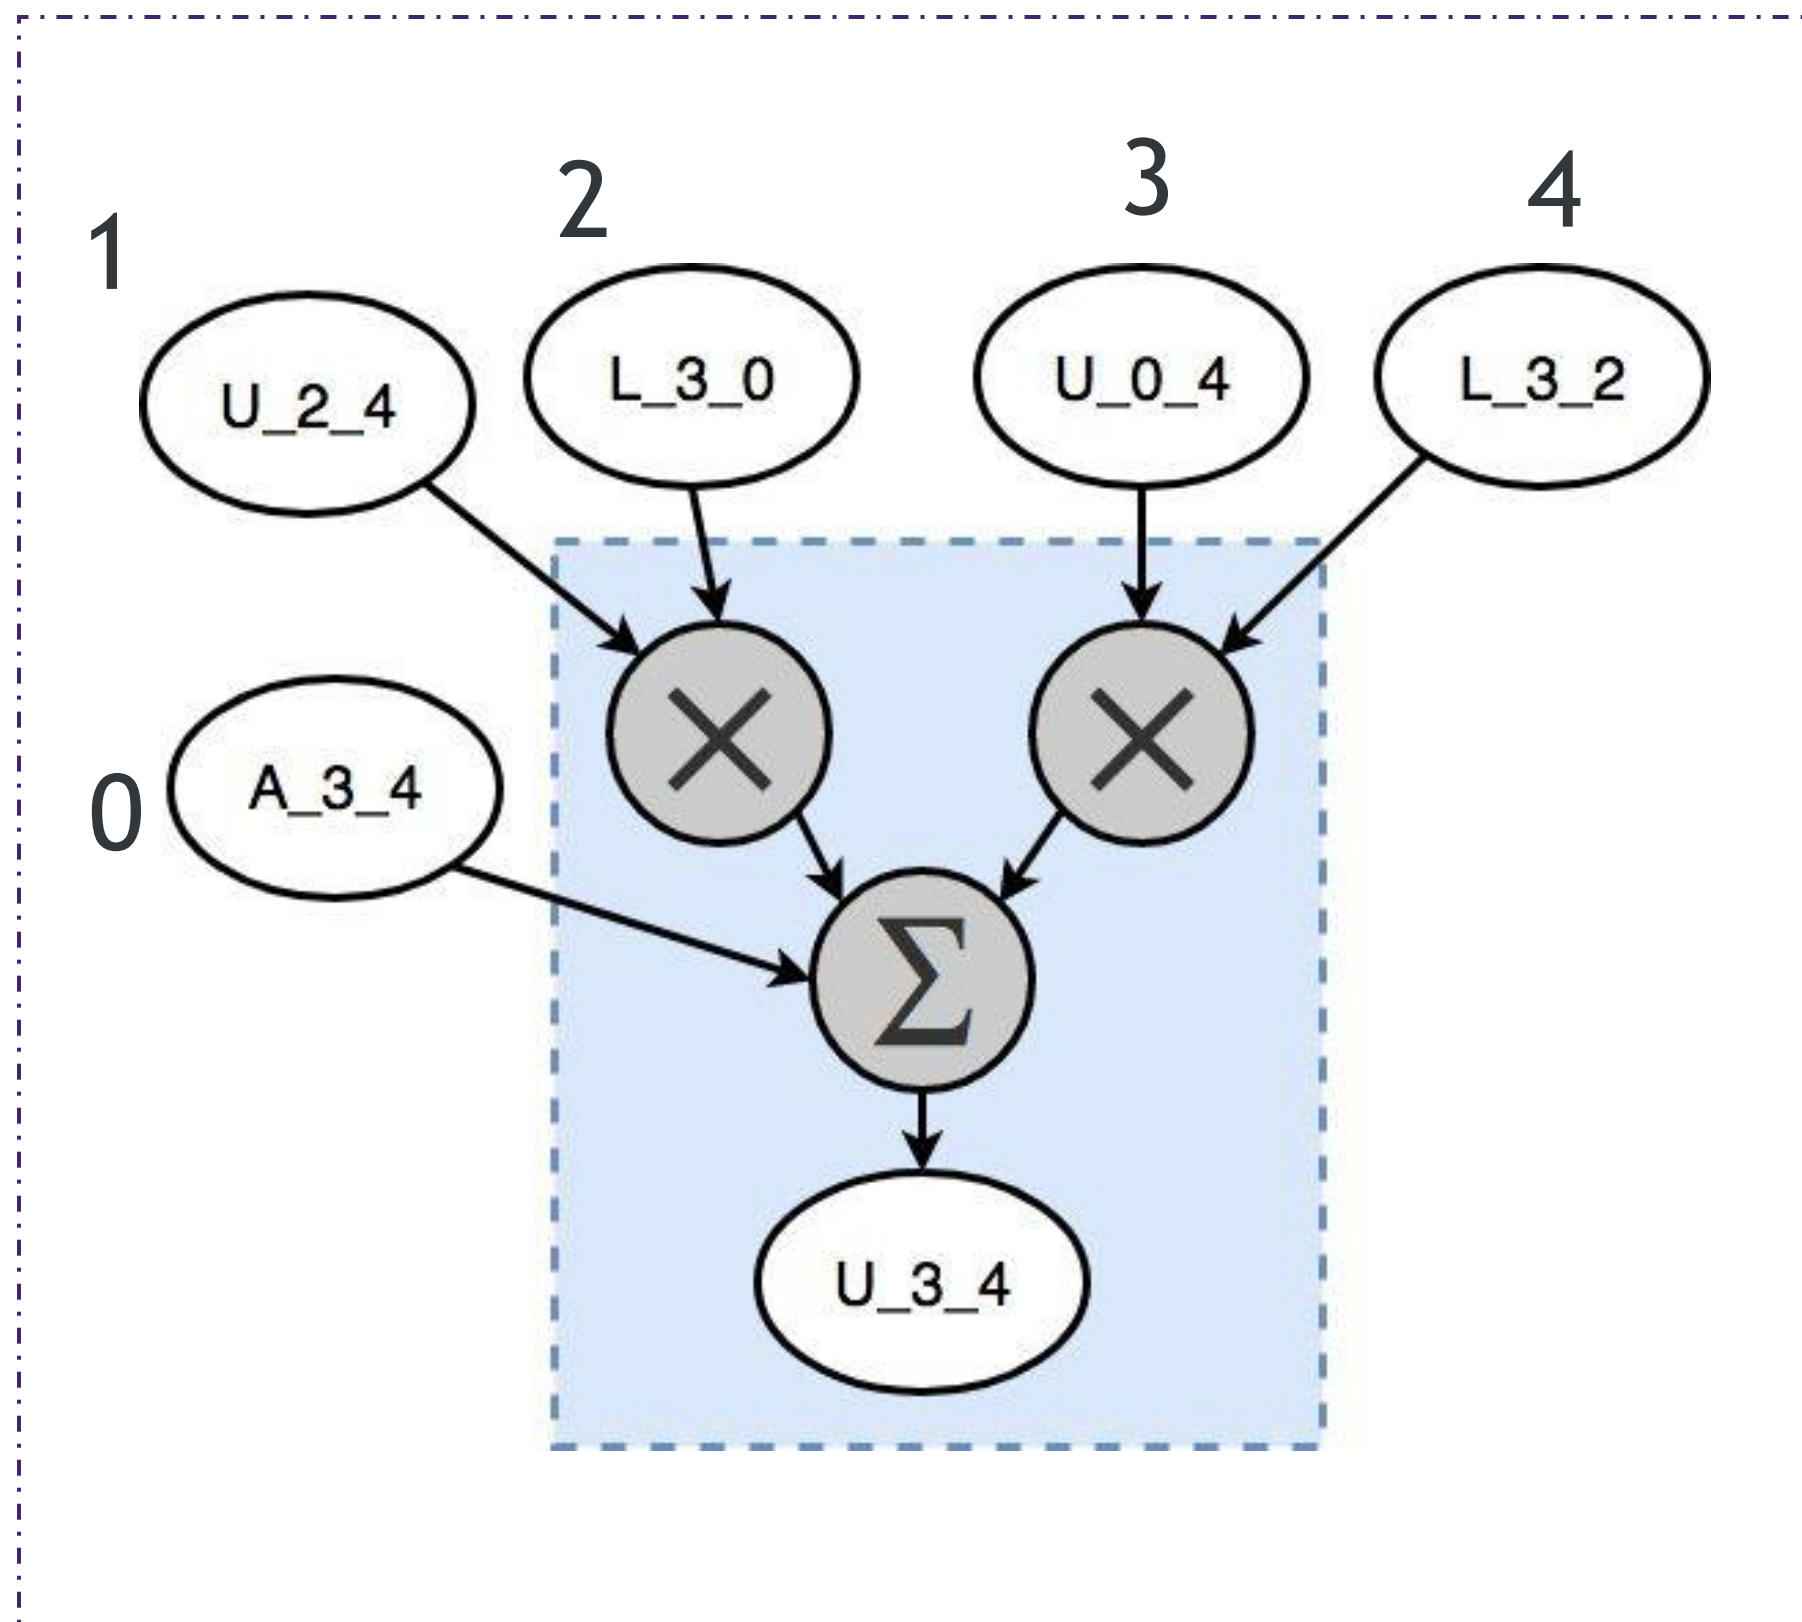




**Step 5:
Column 6**
 $\text{Reach}(3,4,6,7) = \{3,7,4,5,6,9\}$




**Step 6:
Column 7**
 $\text{Rech}(7) = \{7\}$



Computation Task Graph




MAC Super Node

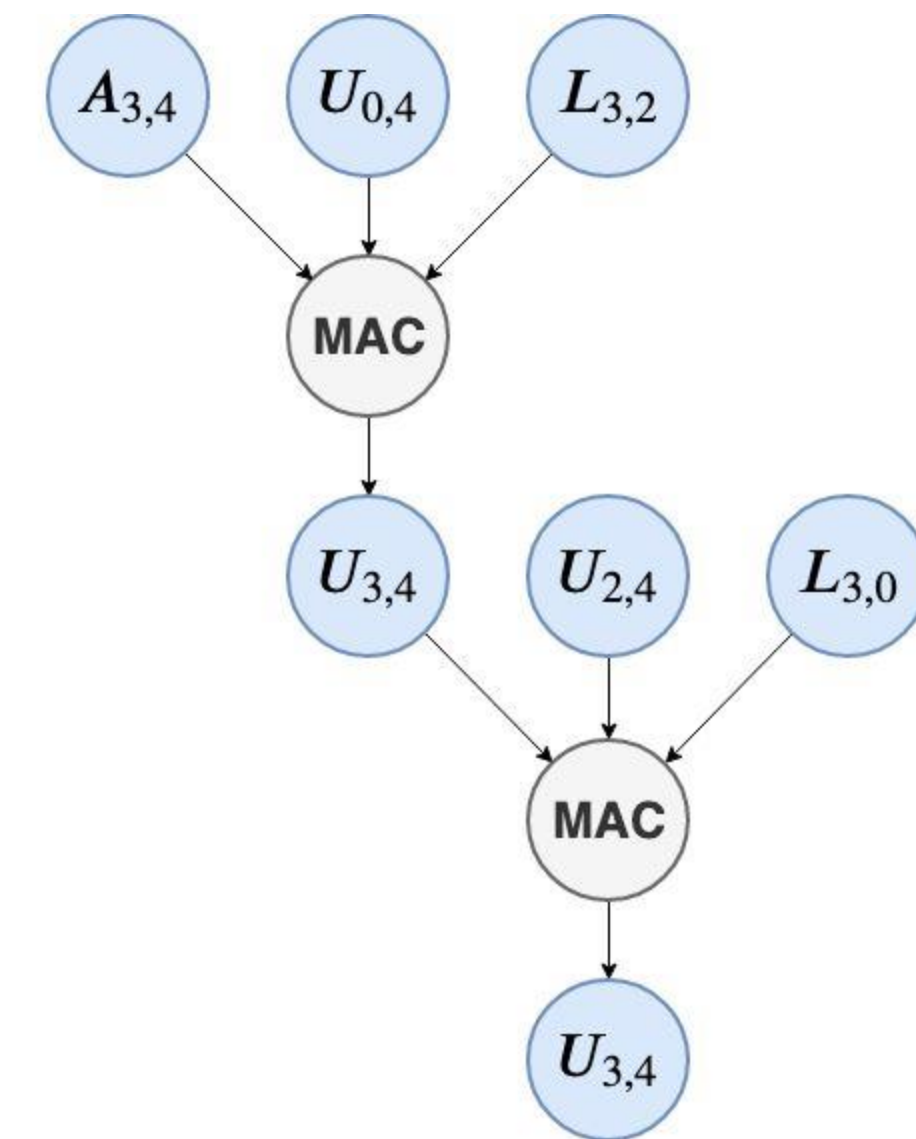


$$U_{3,4} = A_{3,4} - U_{3,4}L_{3,0} - U_{0,4}L_{3,2}$$

Total $n!$ different execution sequences are possible for computing an element with n pairs of multiplicands

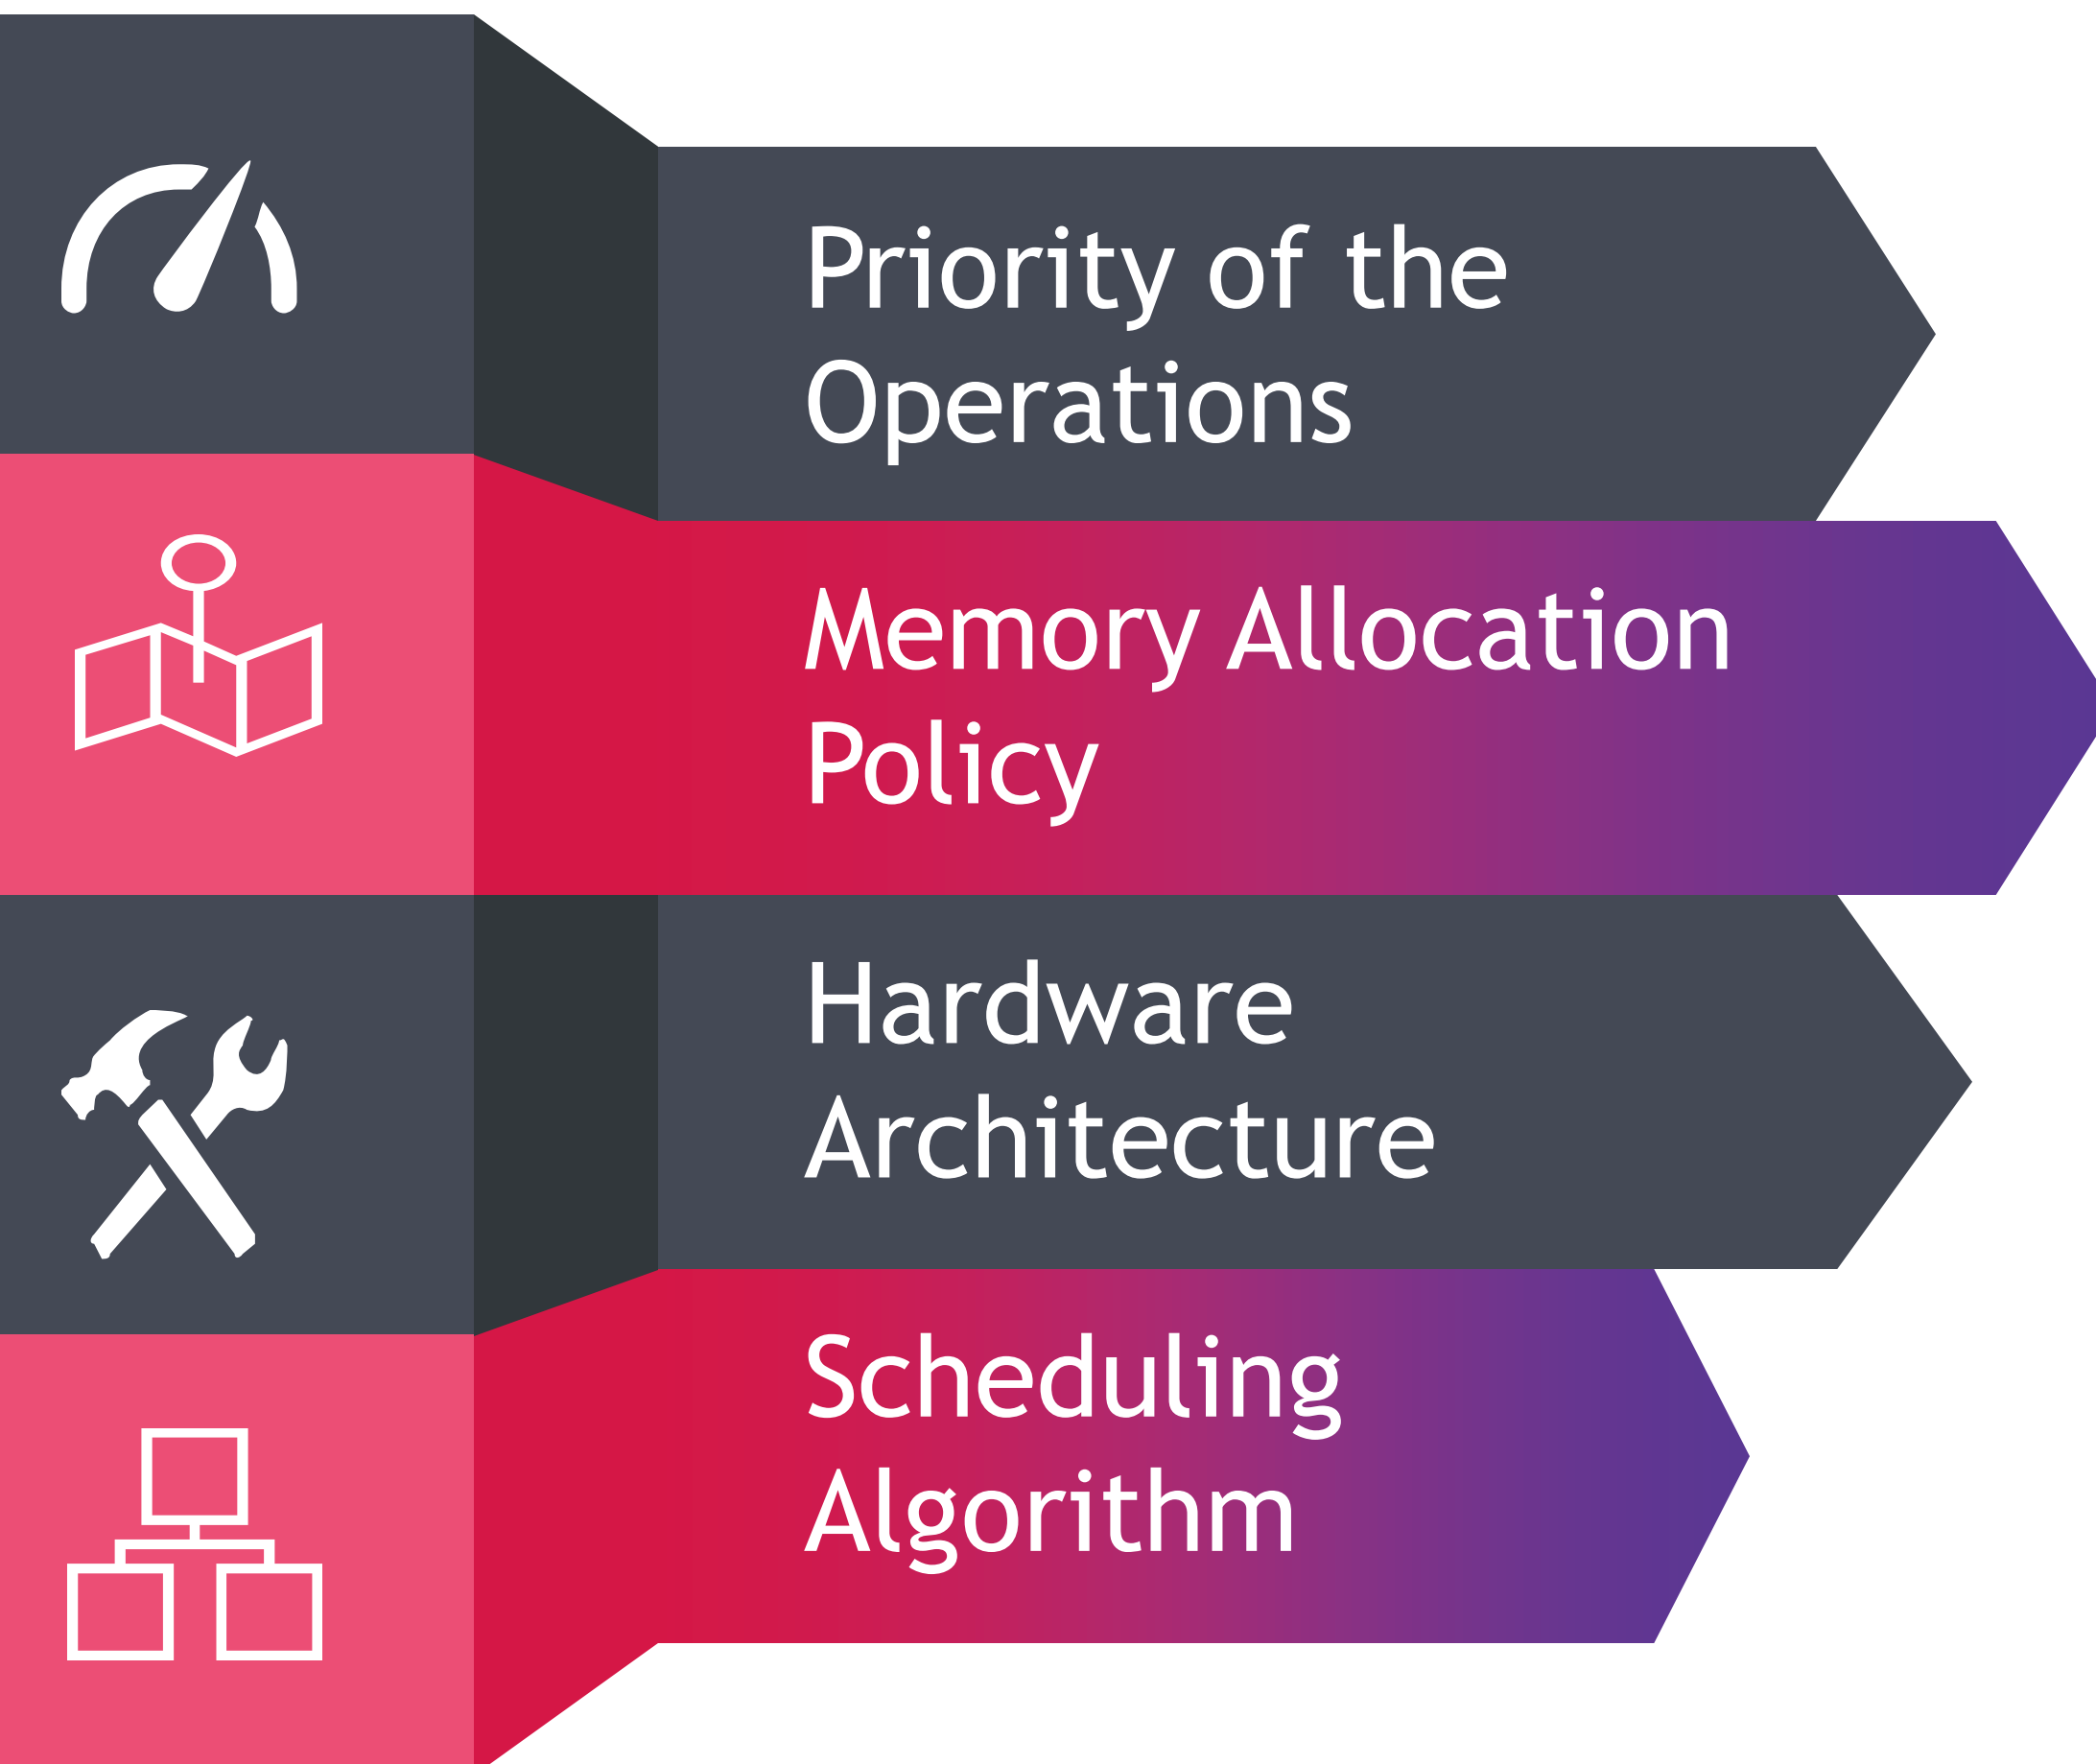


Sequence 1



Sequence 2

Scheduling



Priority of each node is computed recursively using

$$Priority(n) = \sum_{x \in parents(n)} Priority(x) + \sum_{i \in tasks(n)} Delay(i)$$

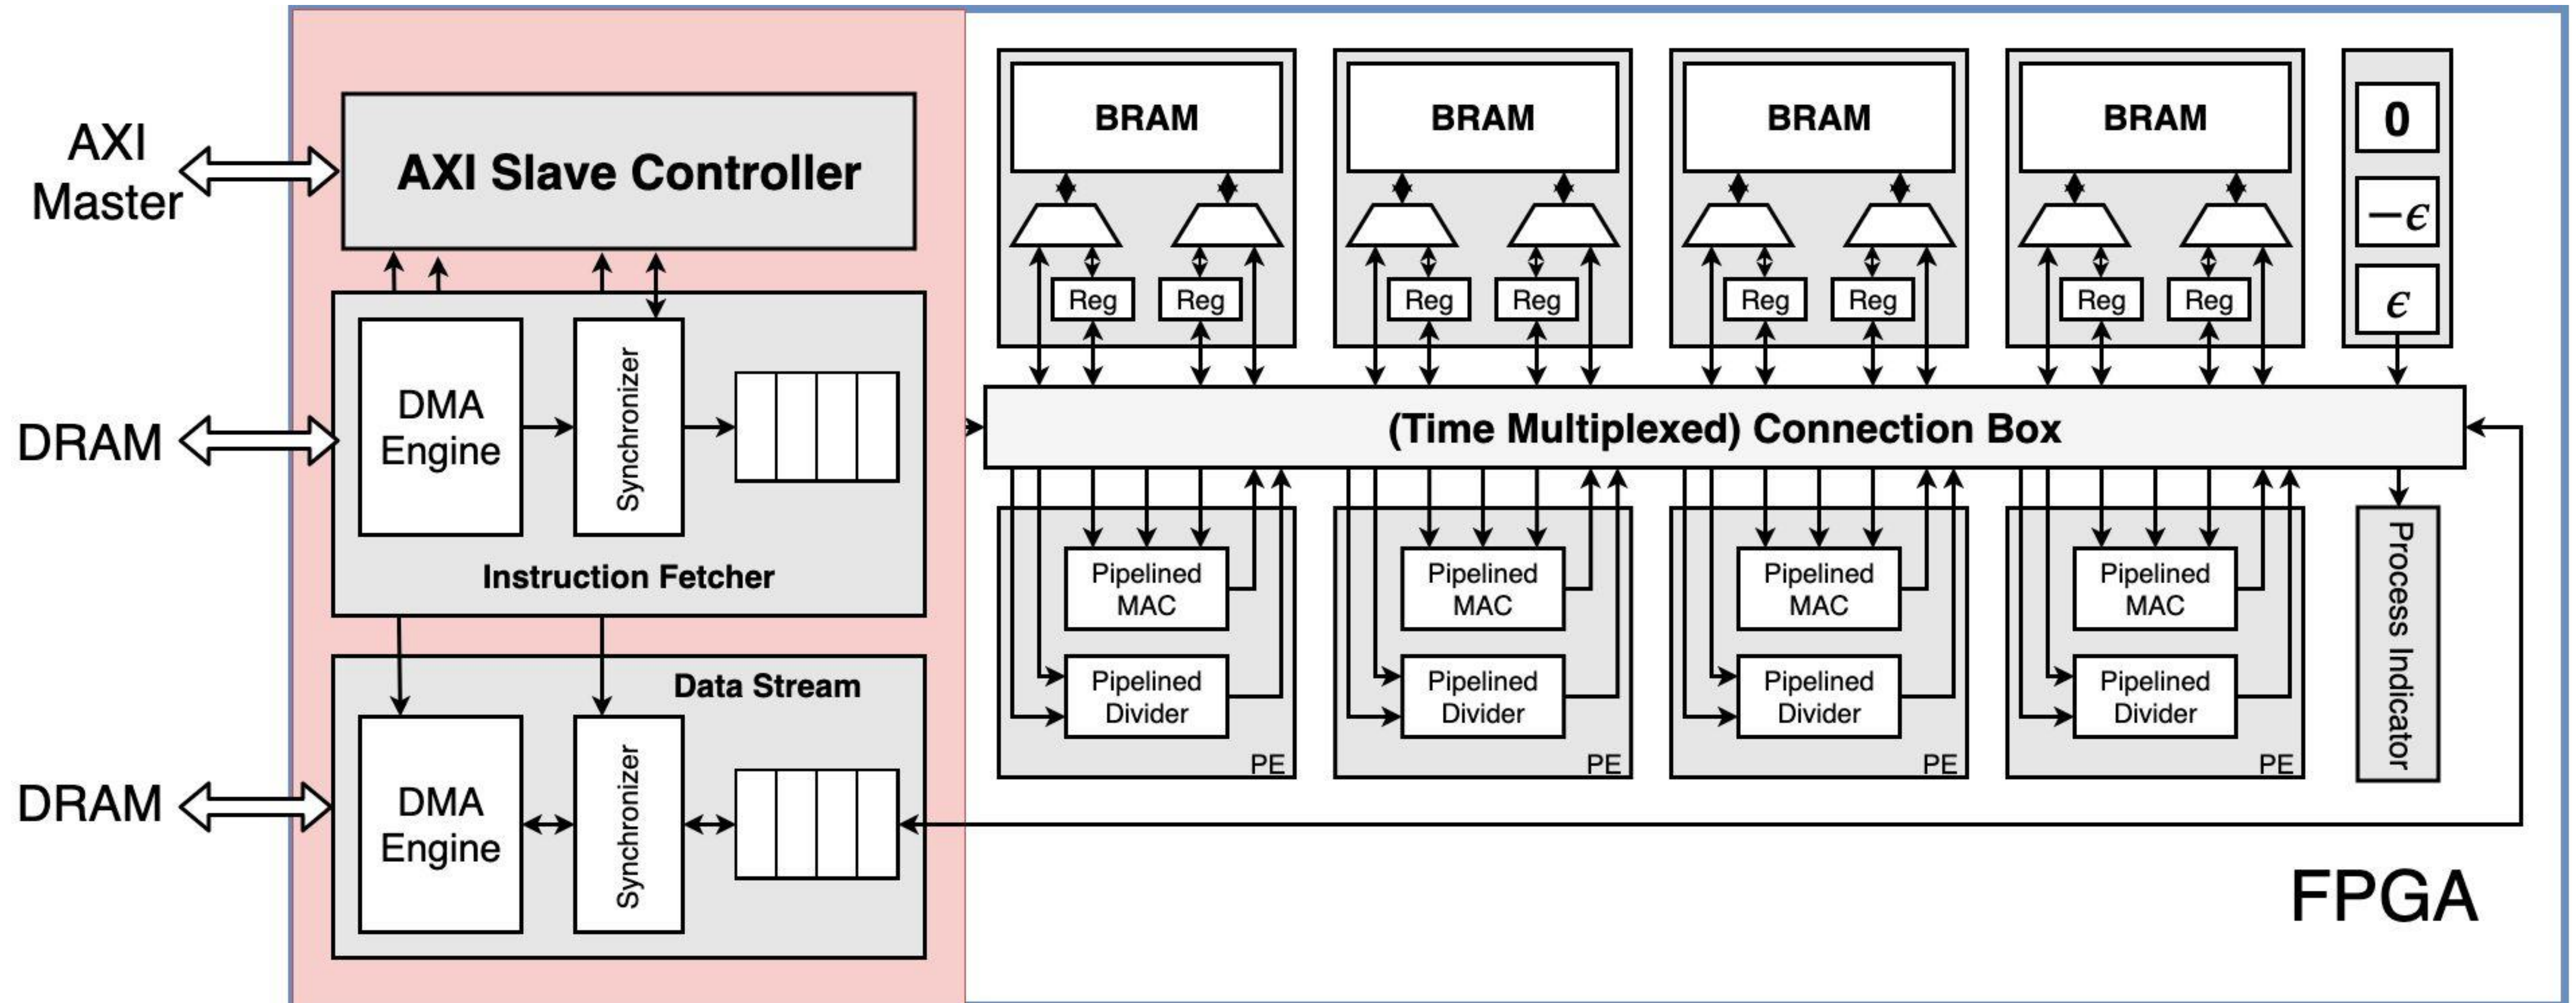
Column wise Memory Allocation

	BRAM 0		BRAM 1		BRAM 2
0	$A_{(0,0)}, U_{(0,0)}$	0	$A_{(1,1)}, U_{(1,1)}$	0	$A_{(0,2)}, U_{(0,2)}$
1	$A_{(2,0)}, L_{(2,0)}$	1	$A_{(3,1)}, L_{(3,1)}$	1	$U_{(2,2)}$
2	$A_{(3,0)}, L_{(3,0)}$	2		2	$L_{(3,2)}$
3		3		3	$A_{(4,2)}, L_{(4,2)}$
4		4		4	

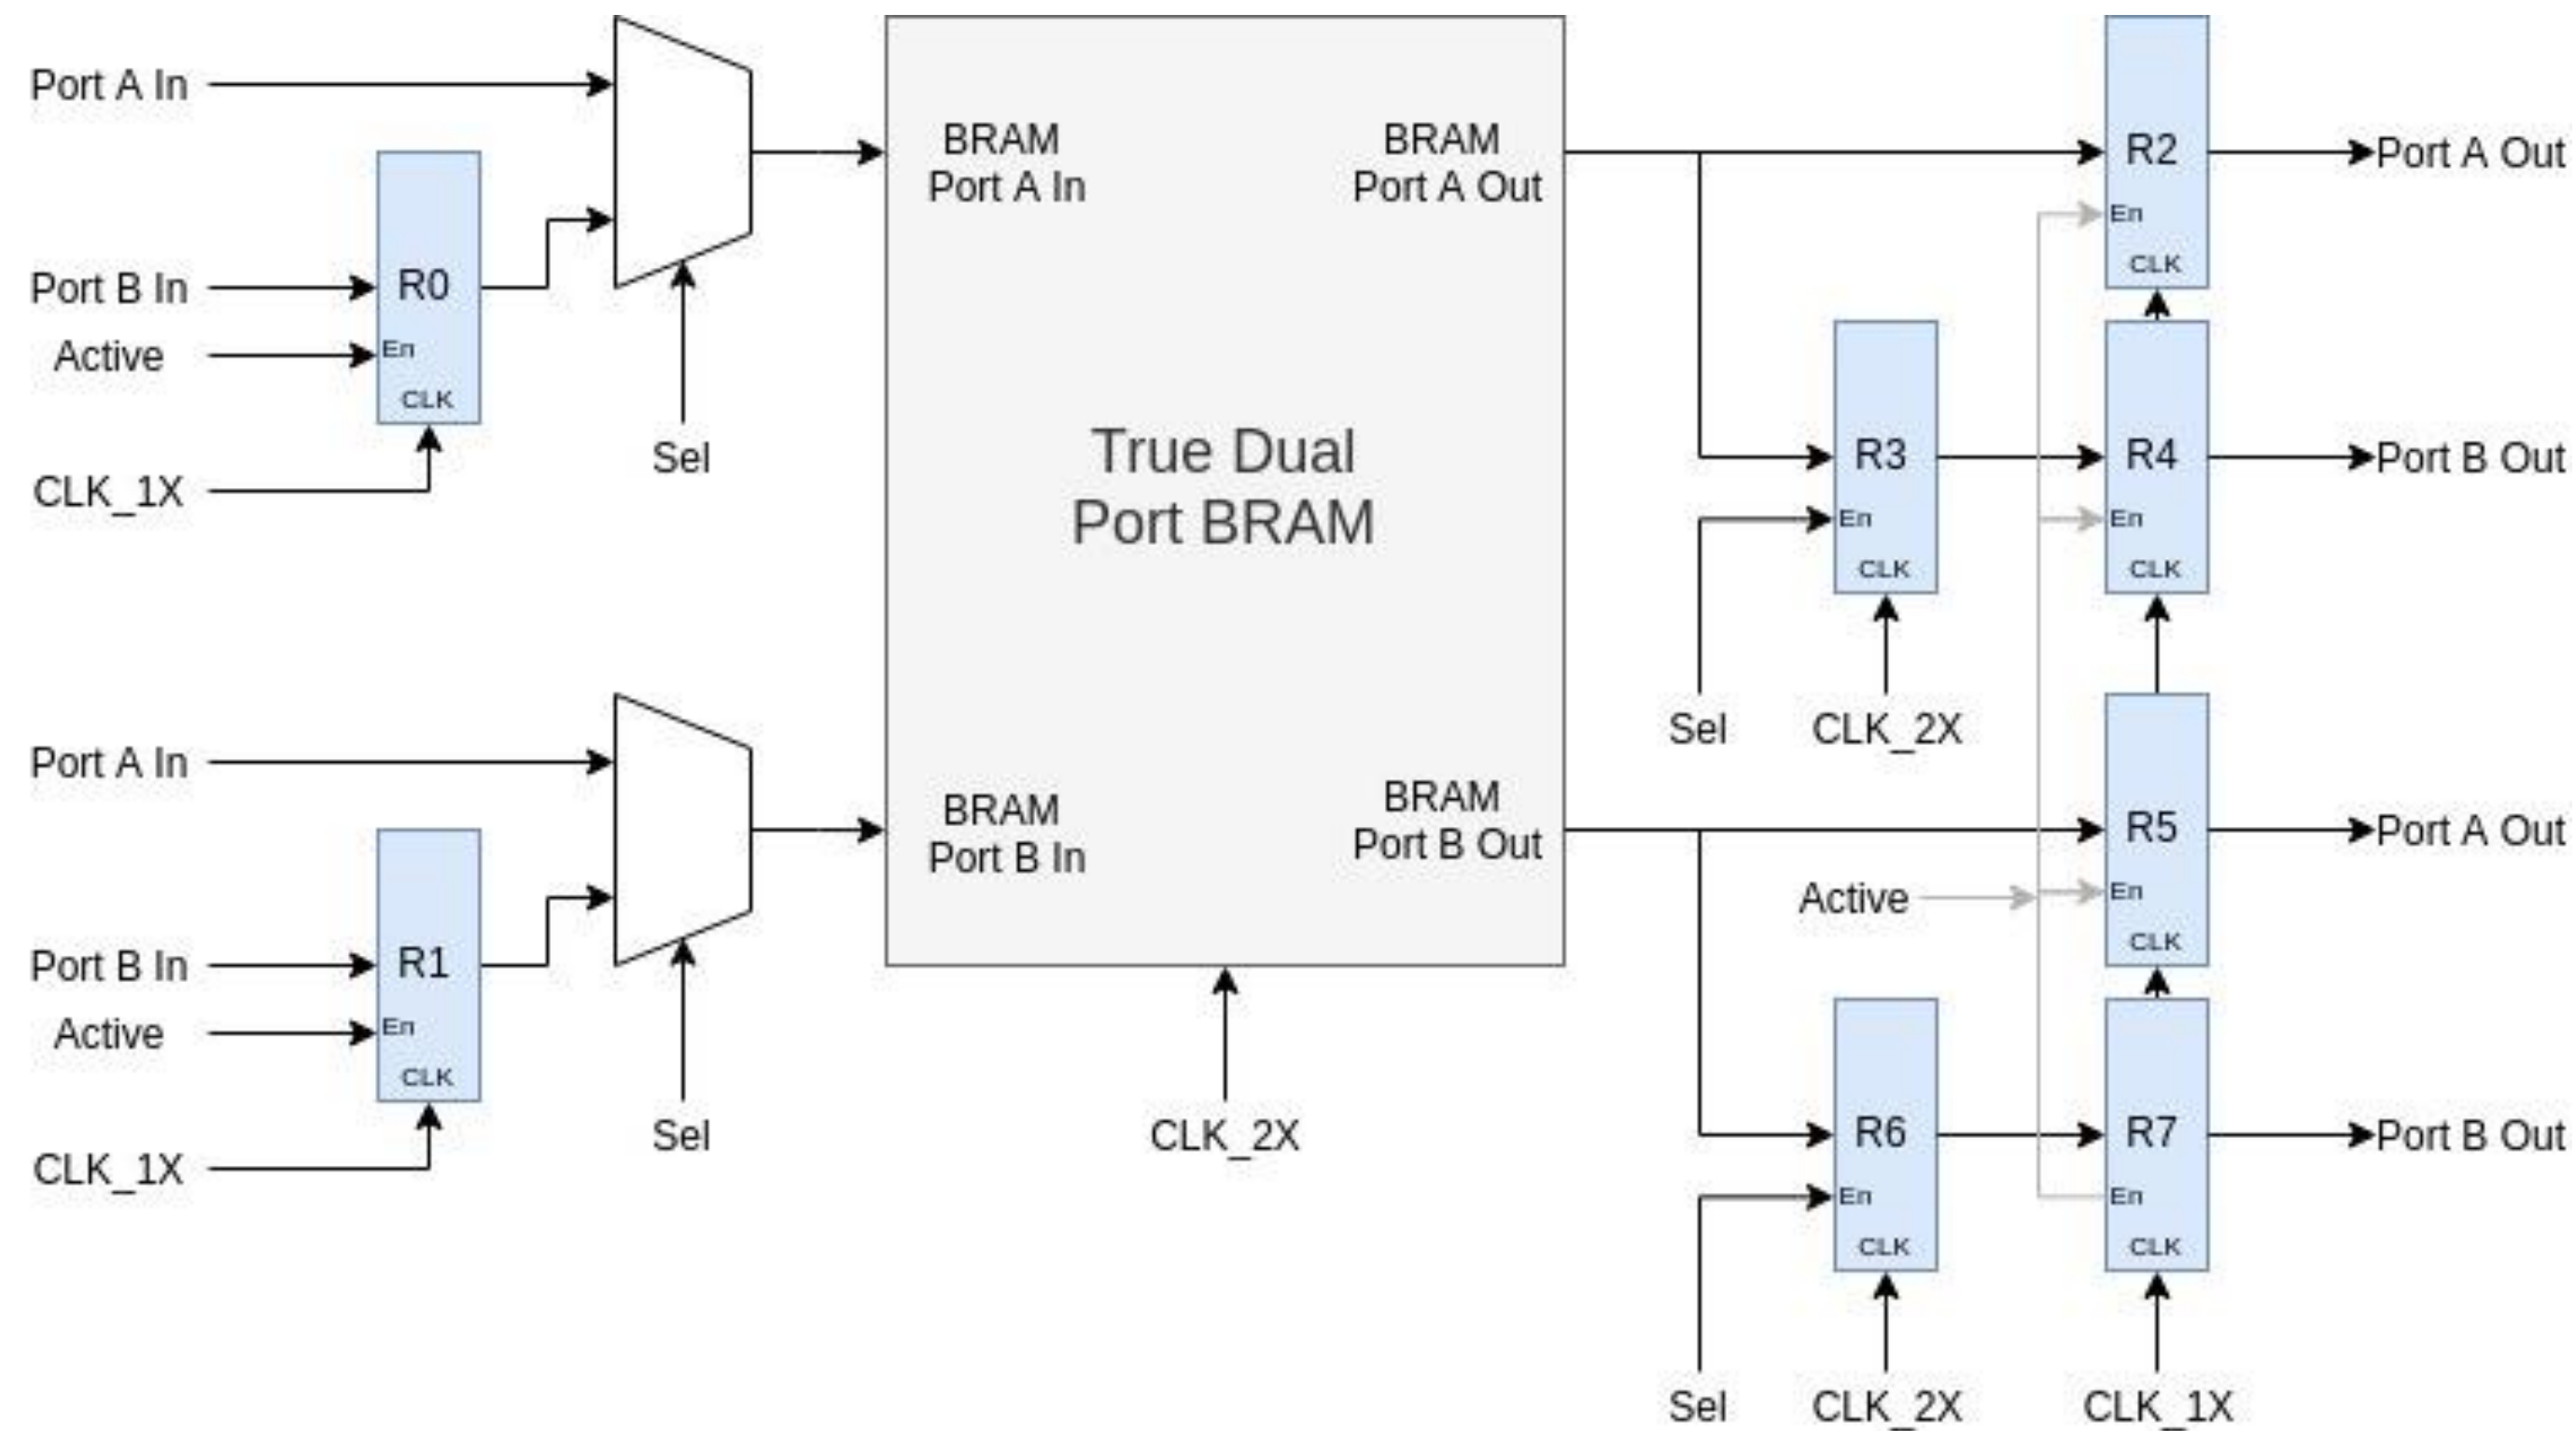
Circular Memory Allocation

	BRAM 0		BRAM 1		BRAM 2
0	$A_{(0,0)}, U_{(0,0)}$	0	$A_{(2,0)}, L_{(2,0)}$	0	$A_{(3,0)}, L_{(3,0)}$
1	$A_{(1,1)}, U_{(1,1)}$	1	$A_{(3,1)}, L_{(3,1)}$	1	$A_{(0,2)}, U_{(0,2)}$
2	$U_{(2,2)}$	2	$L_{(3,2)}$	2	$A_{(4,2)}, L_{(4,2)}$
3		3		3	
4		4		4	

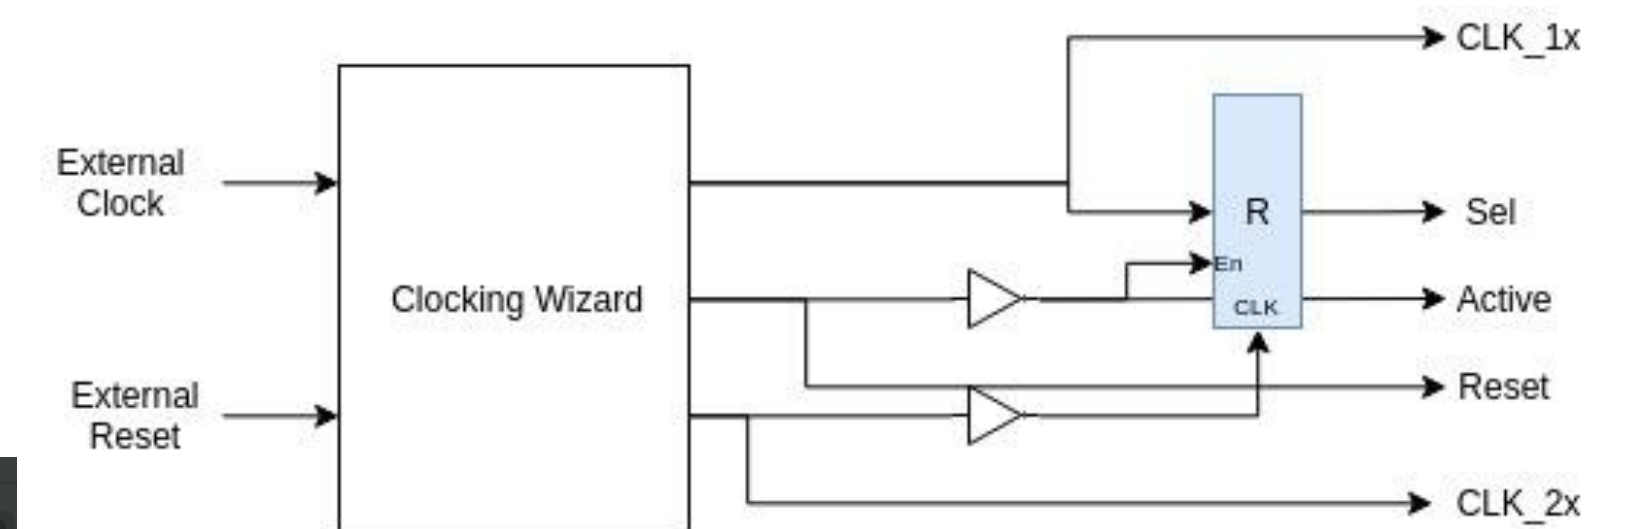
Hardware Architecture



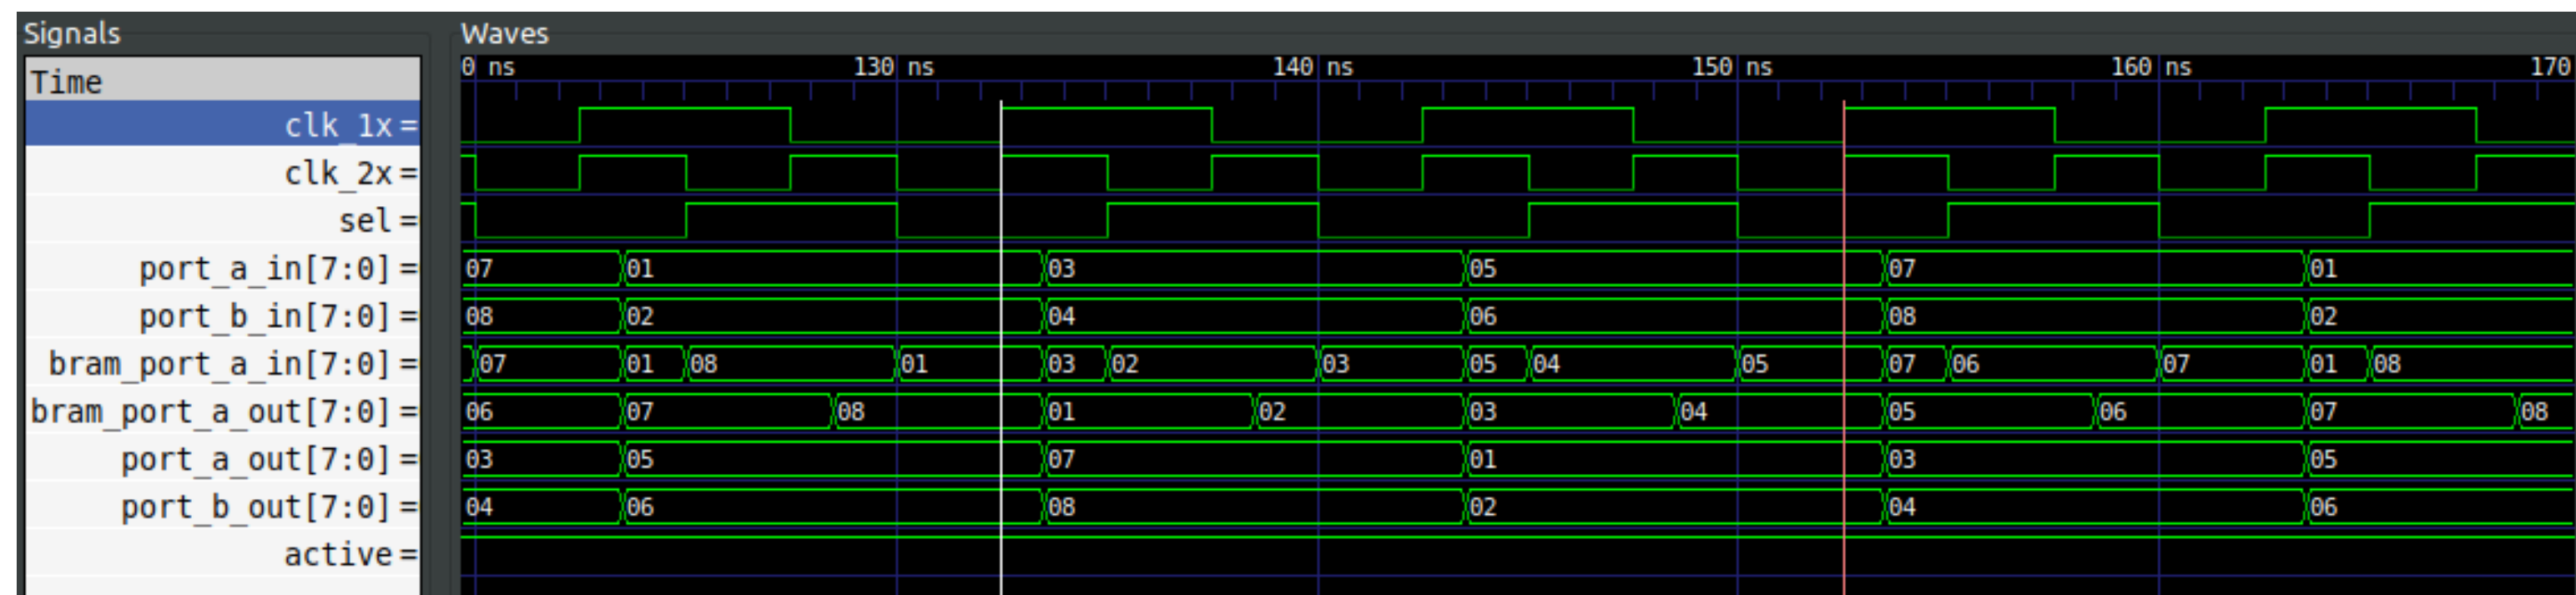
Quad Port BRAM



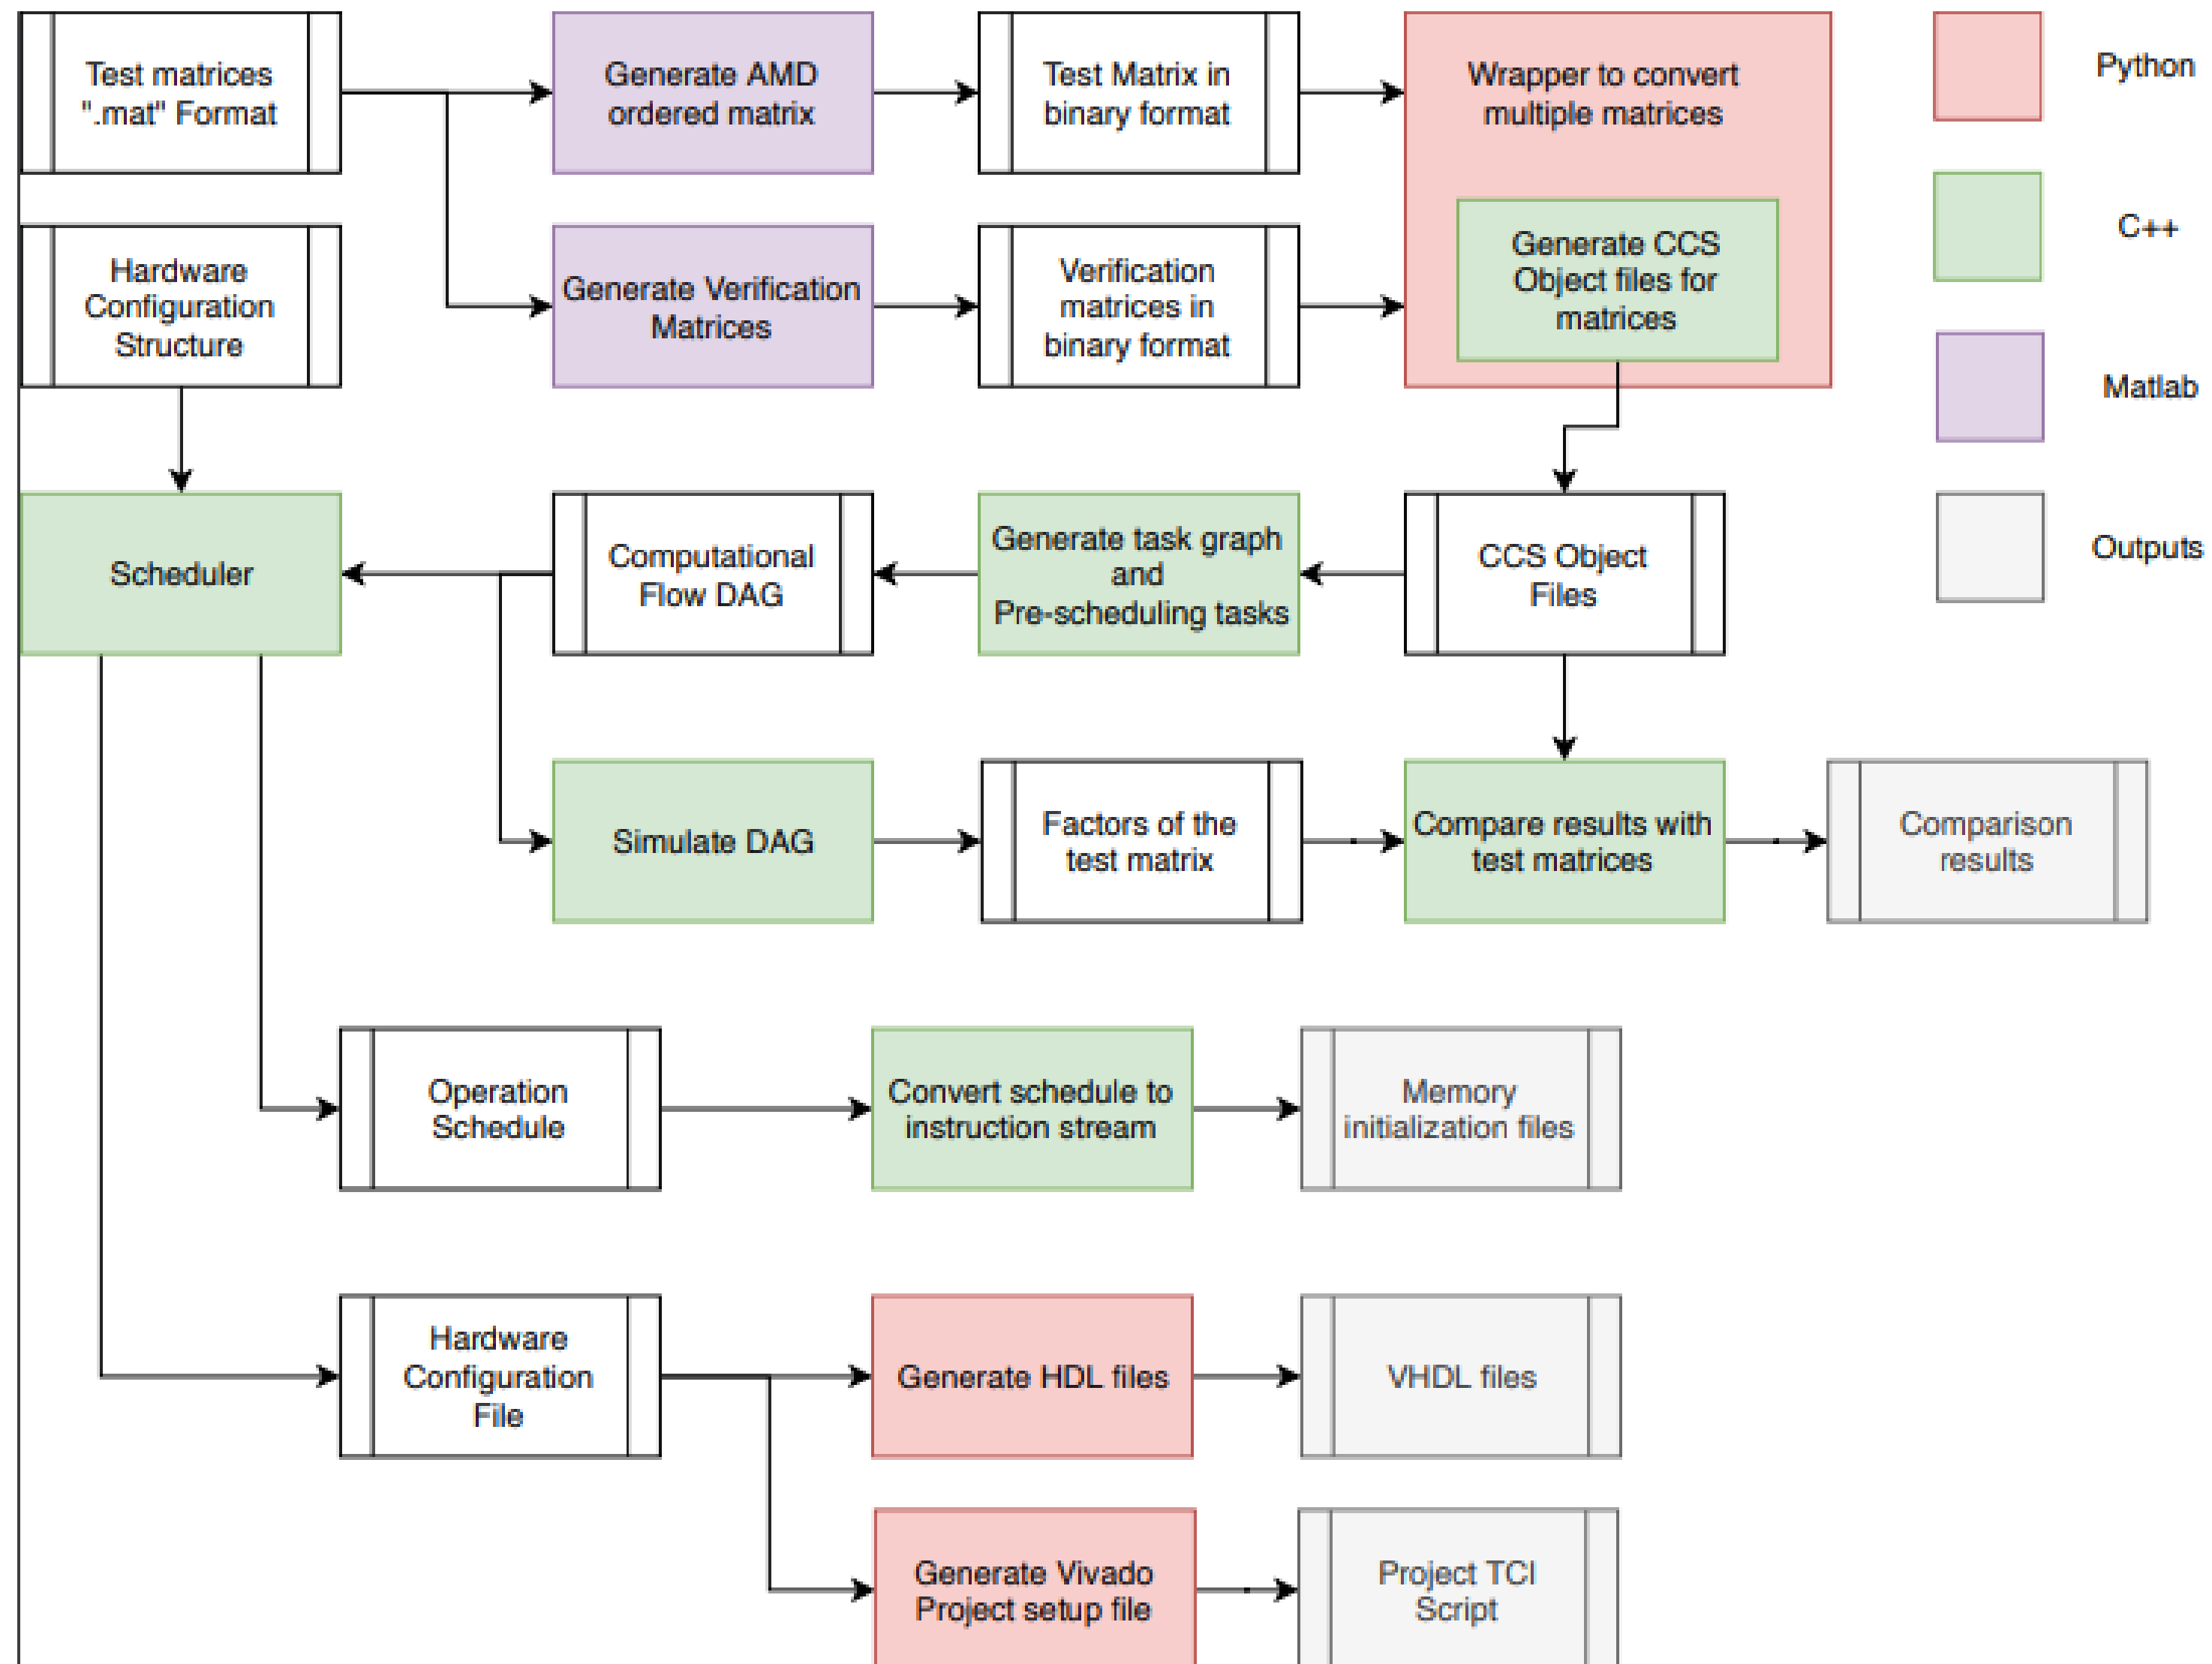
BRAMs can operate at the clock frequencies up to 400 MHz, which is at least two time faster than the processing elements



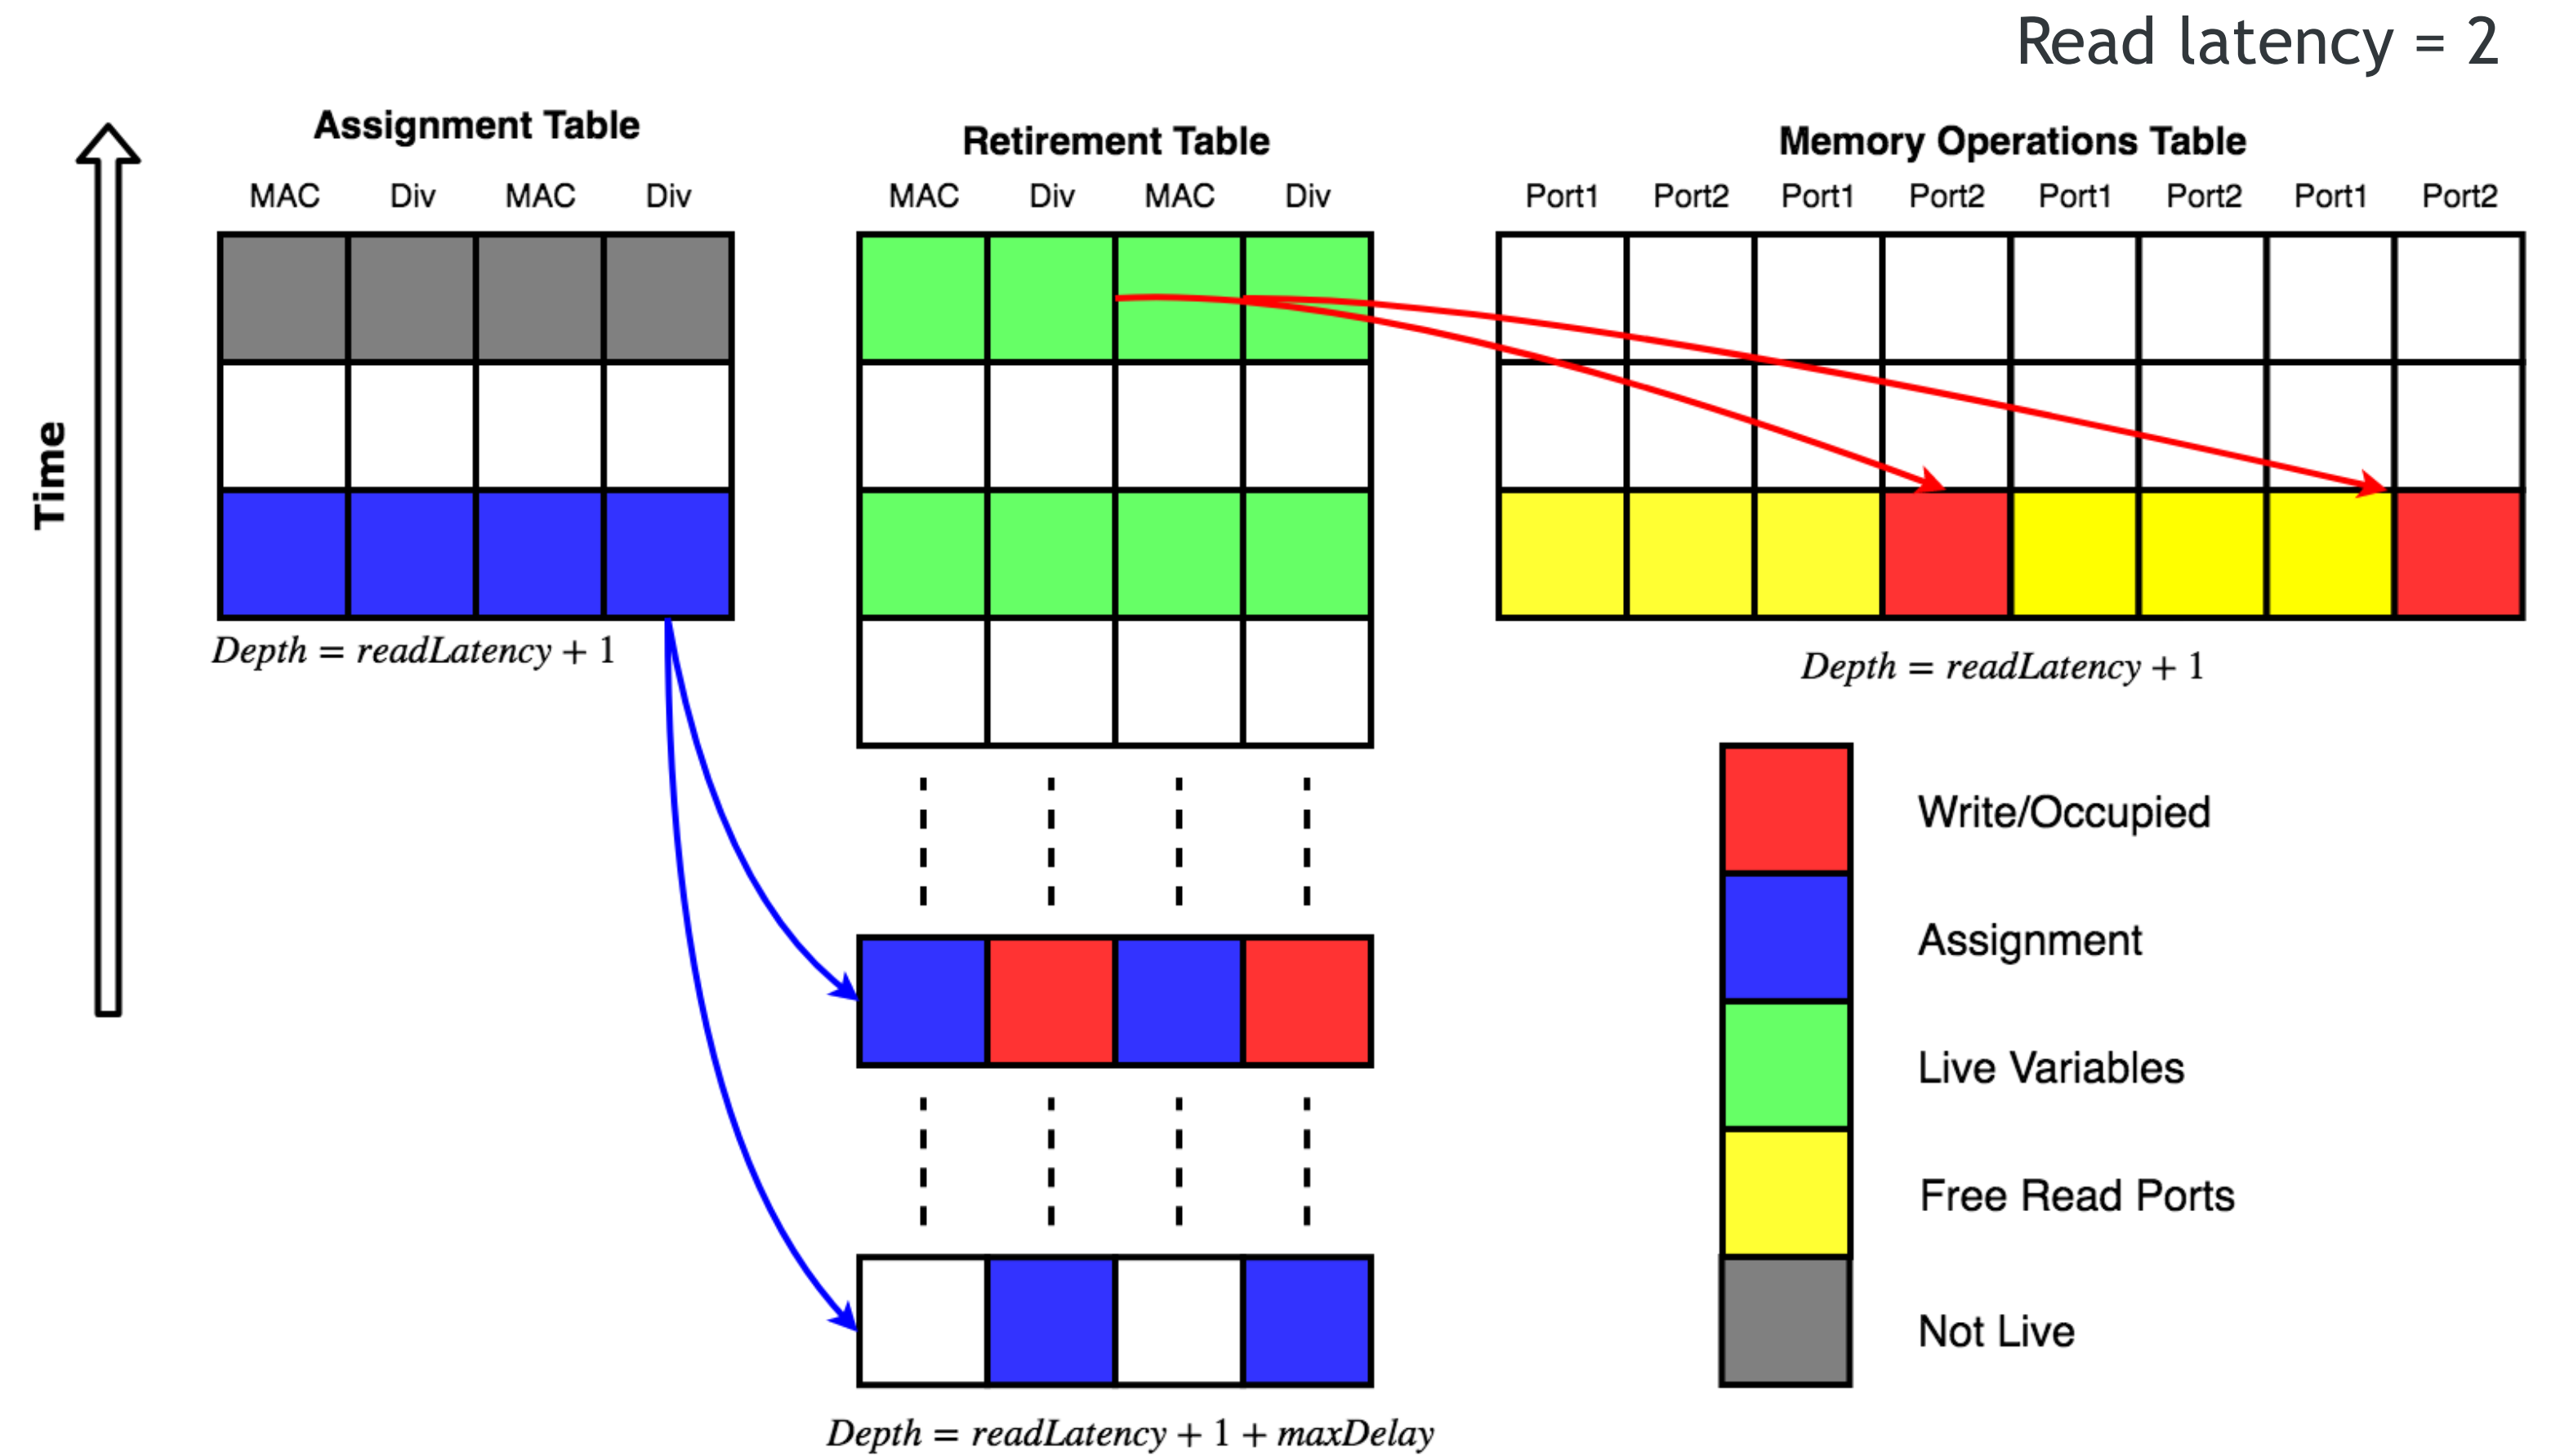
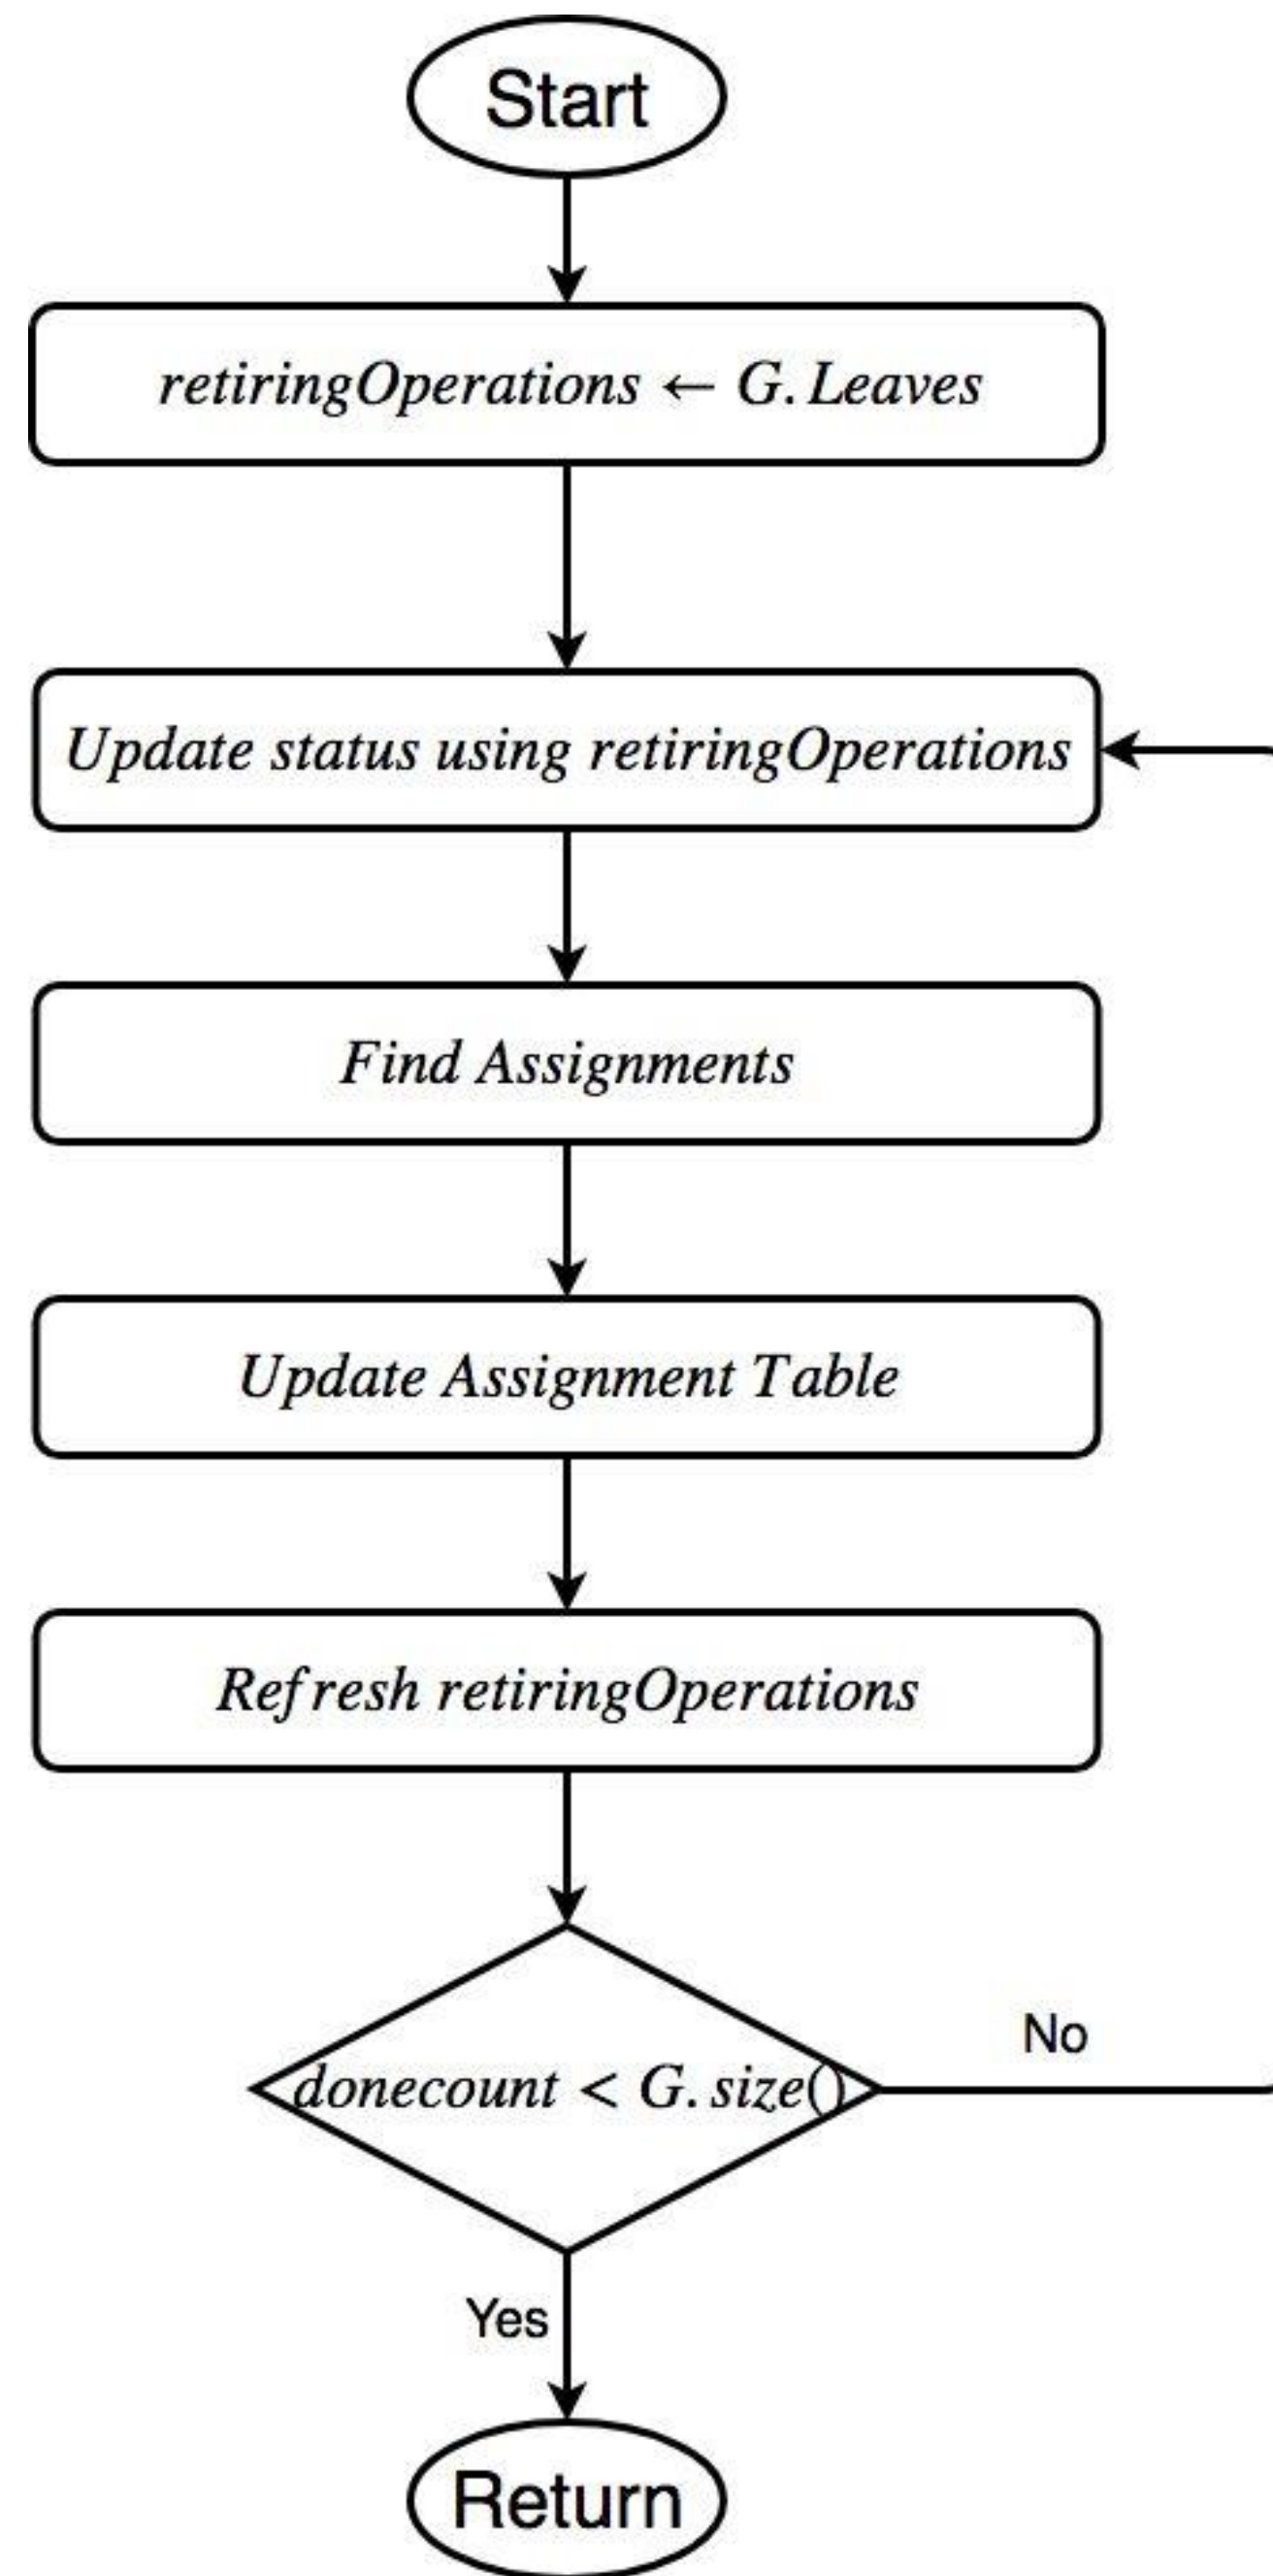
Clock and Reset Generation



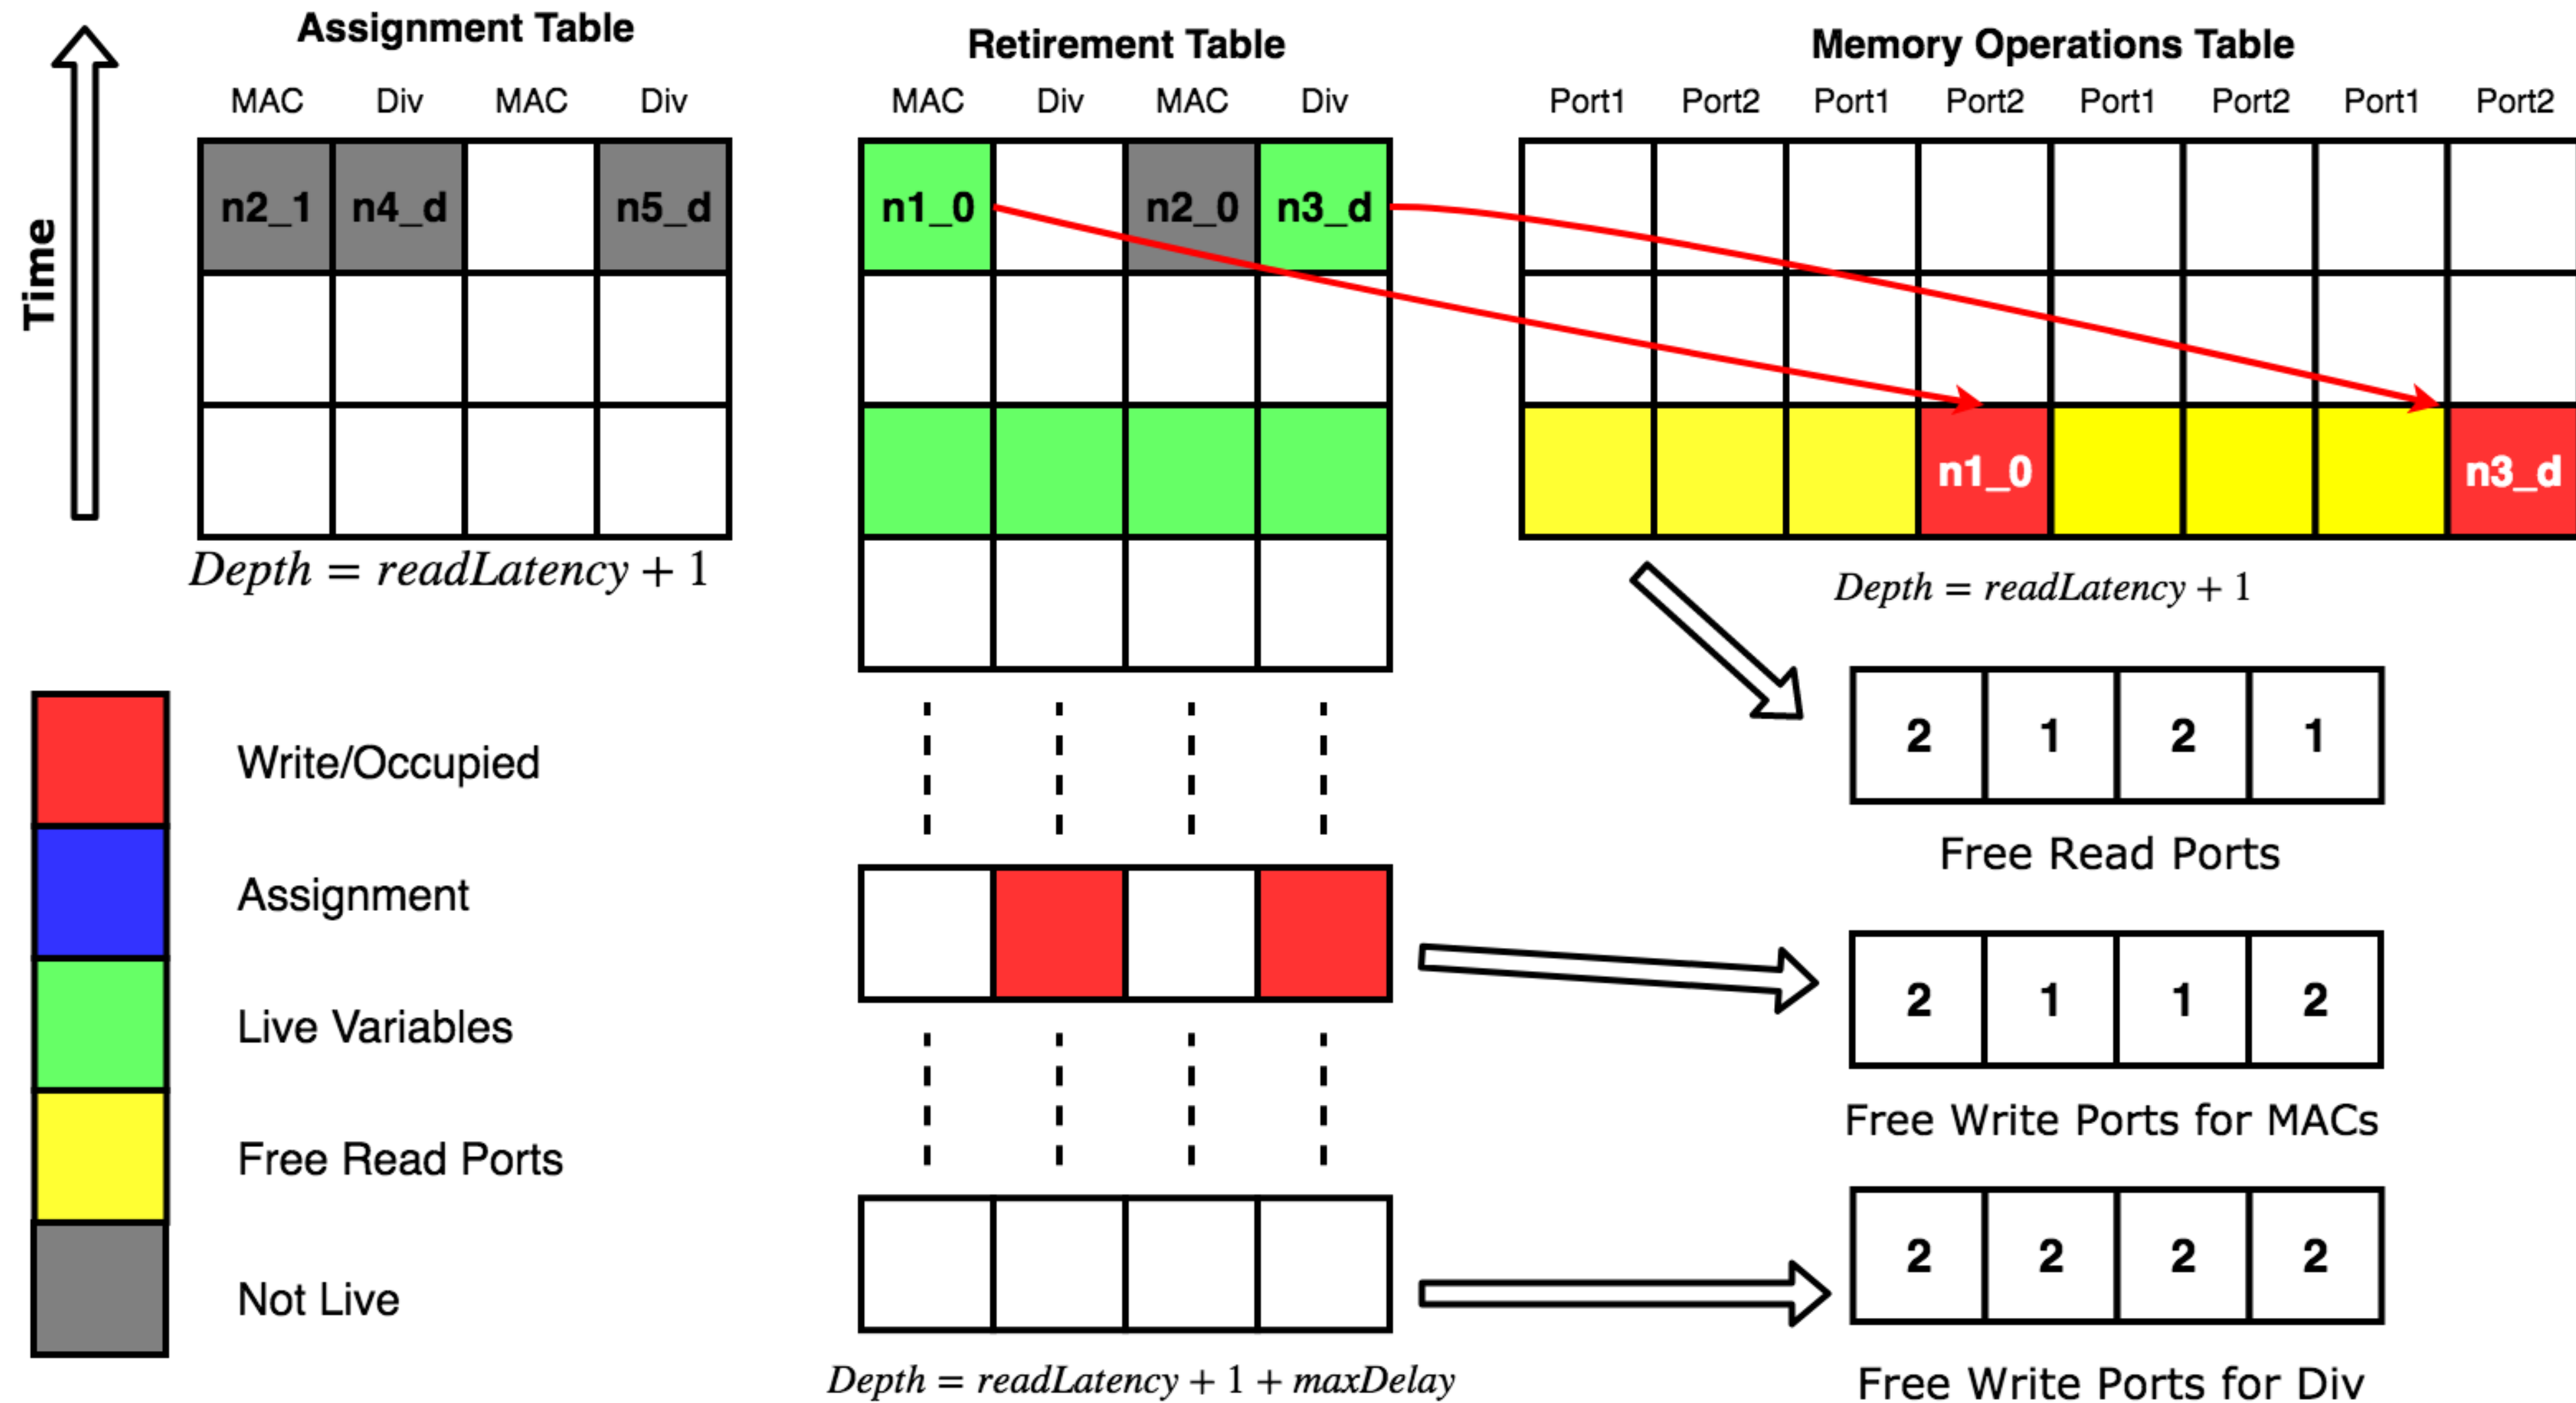
Overall Control Flow



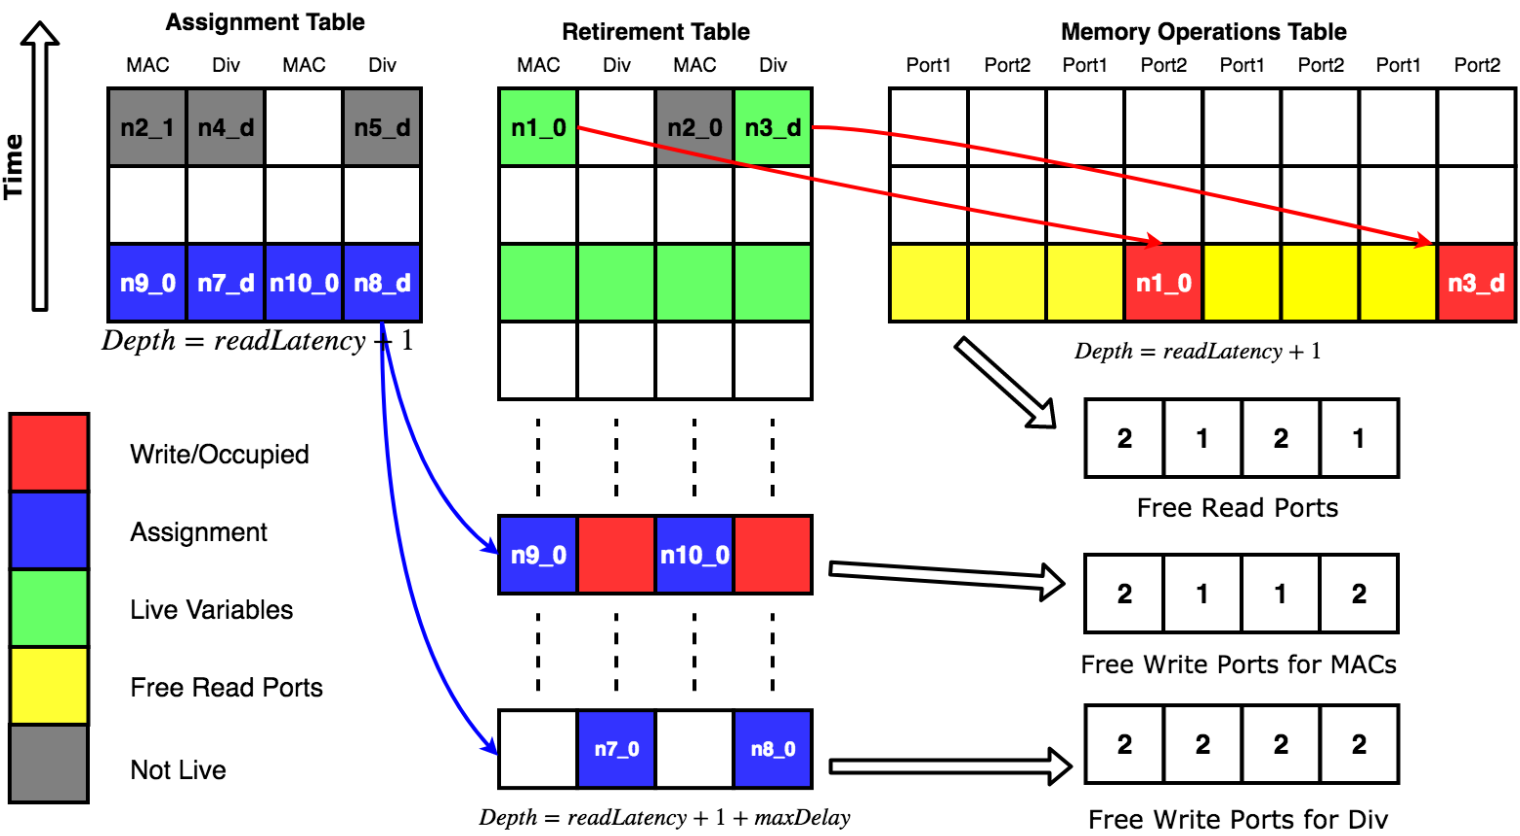
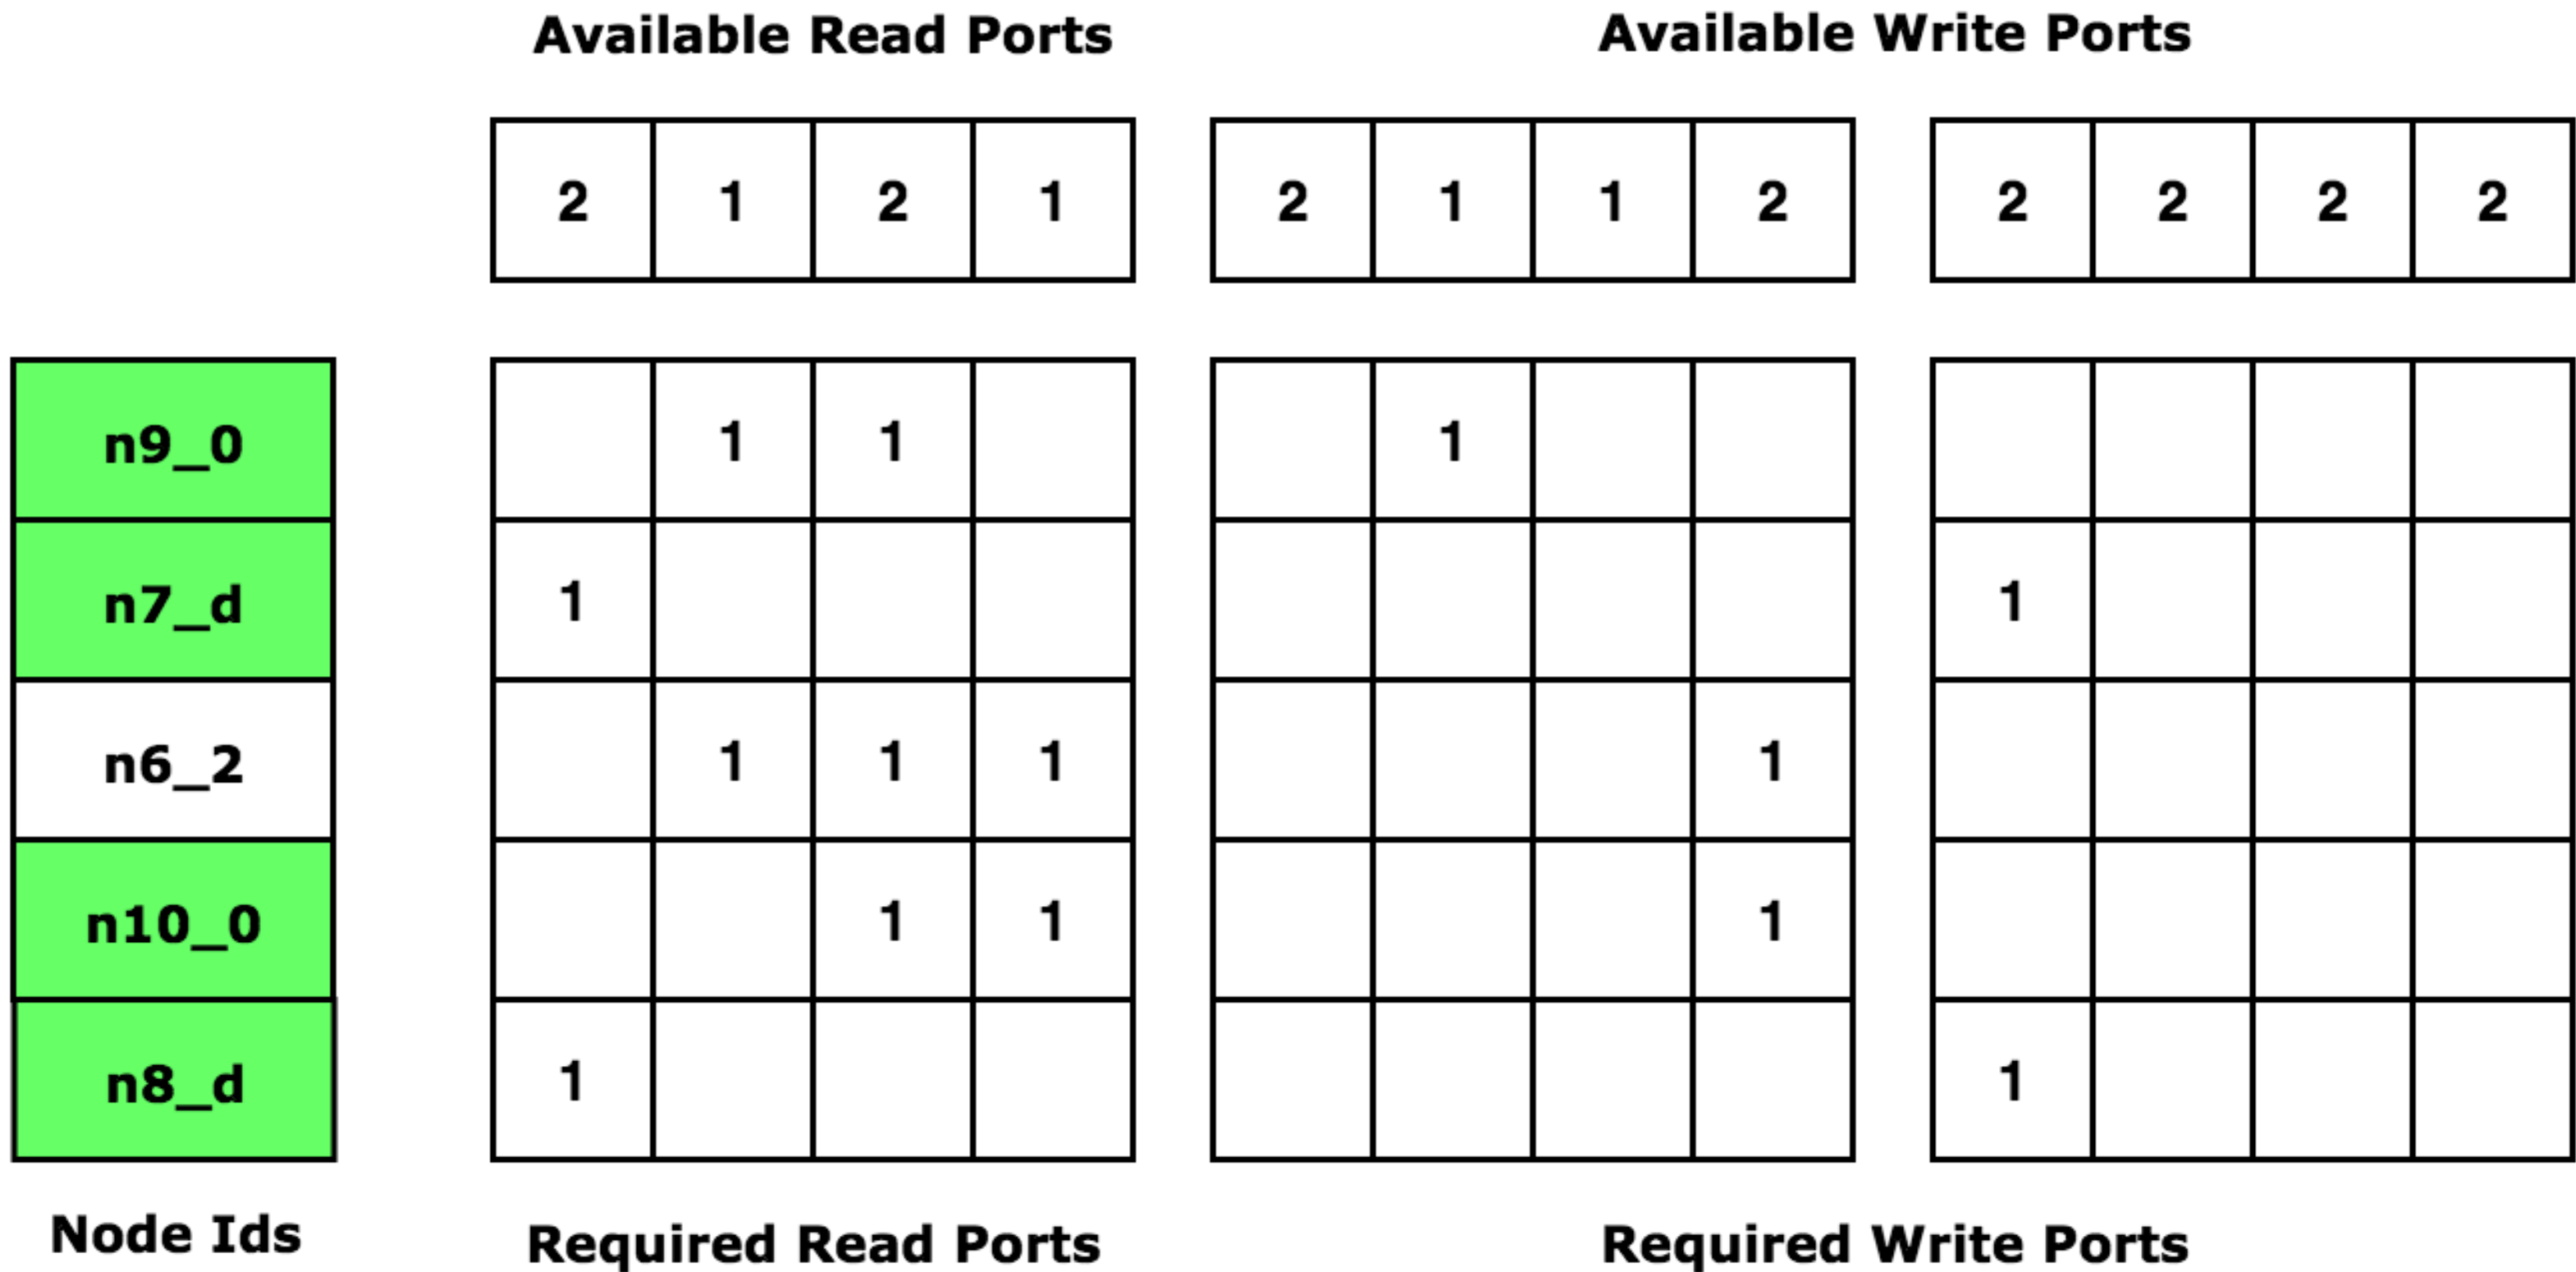
Scheduling



Finding Assignments

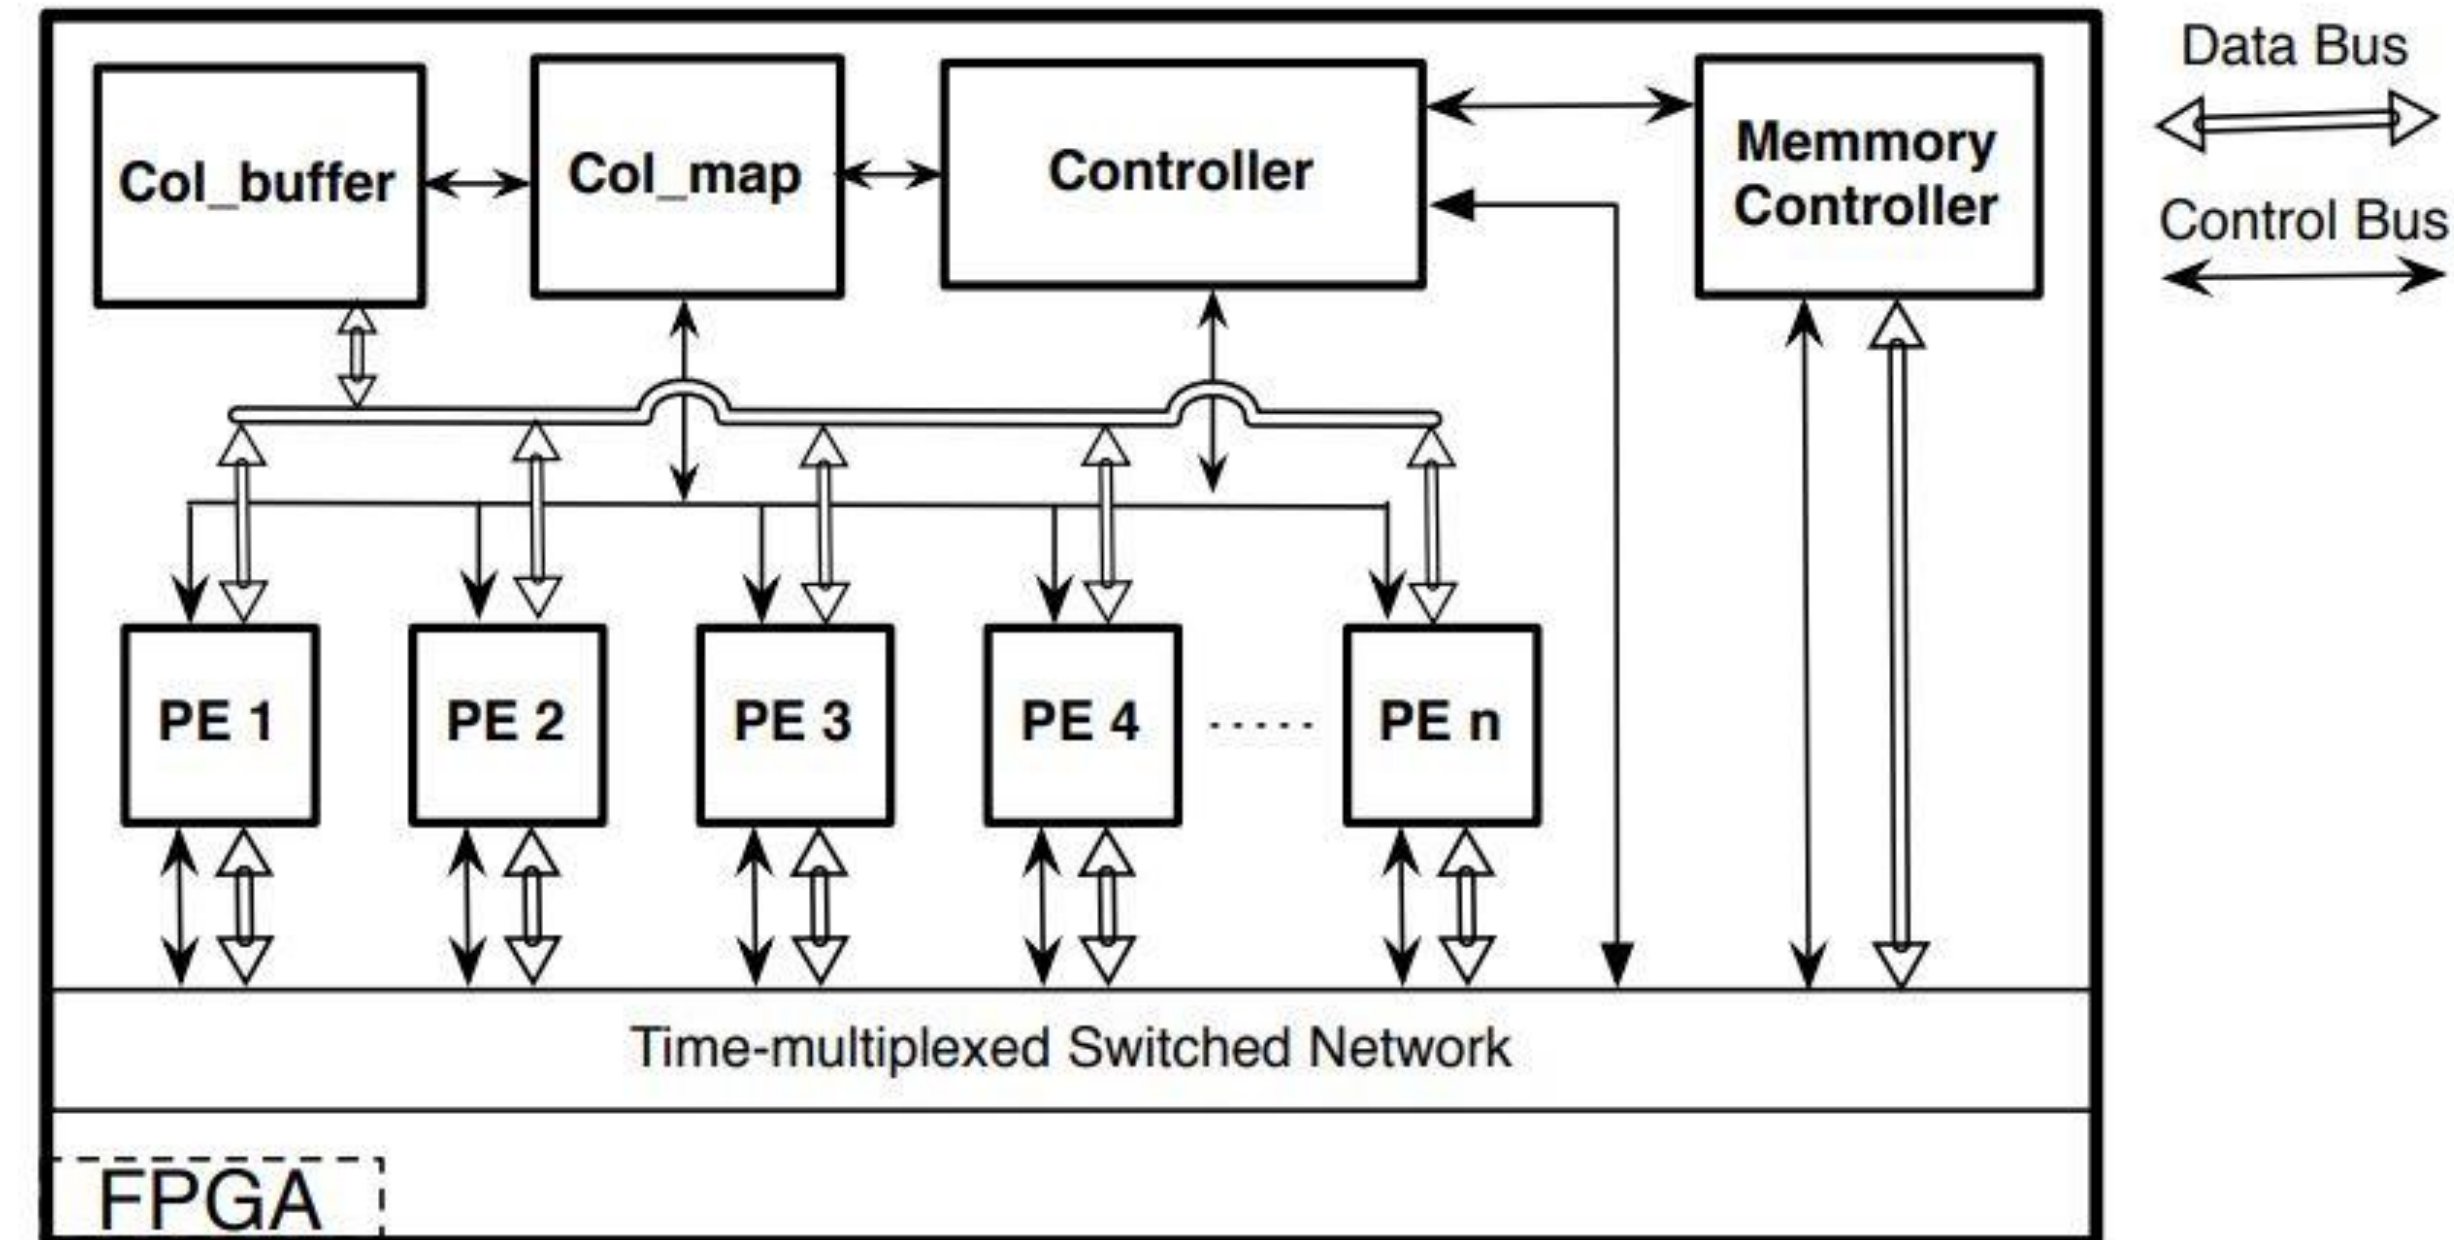
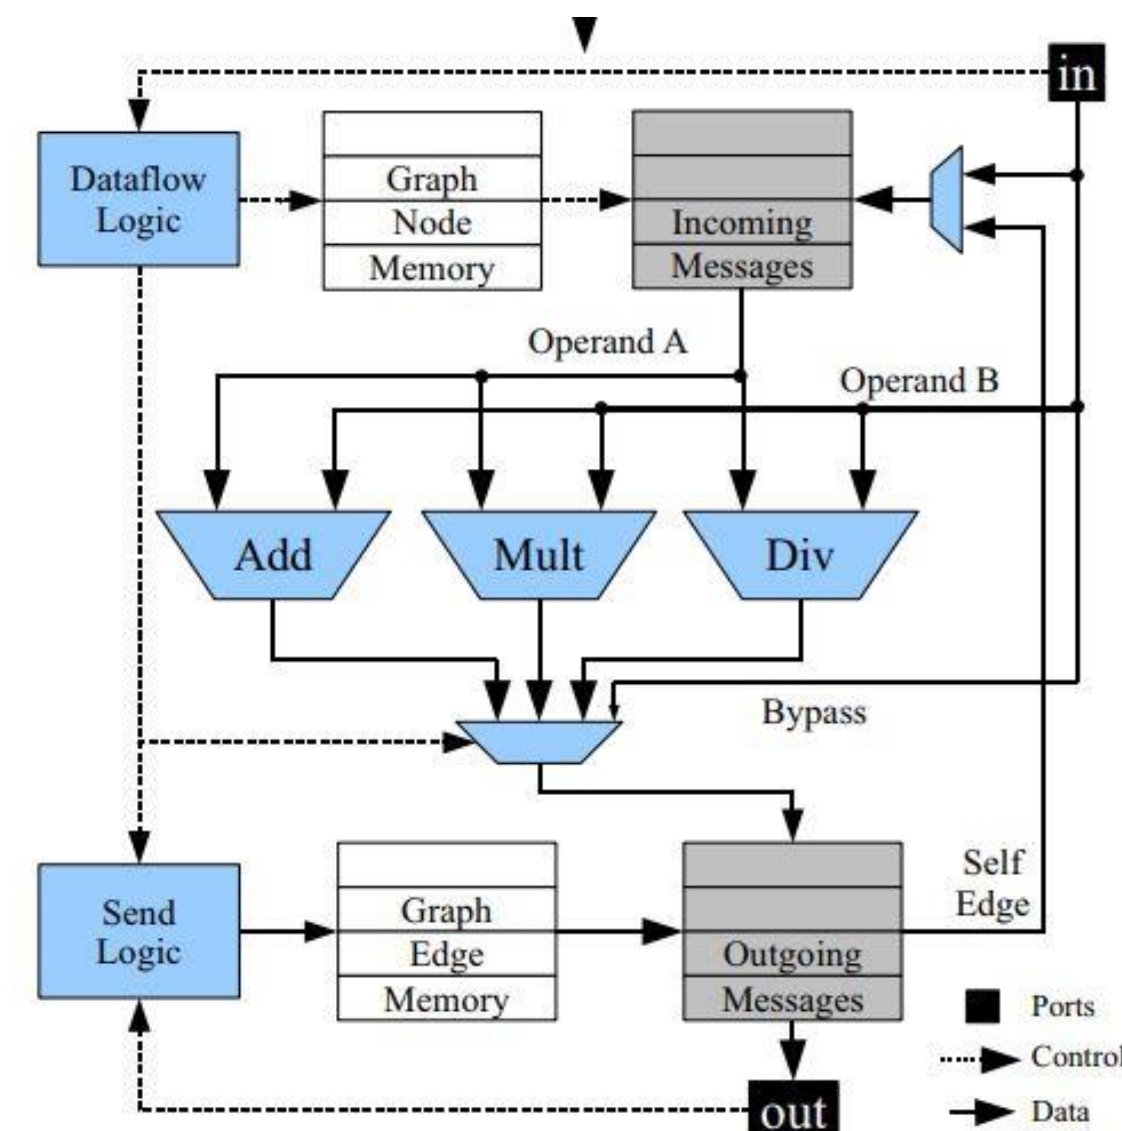
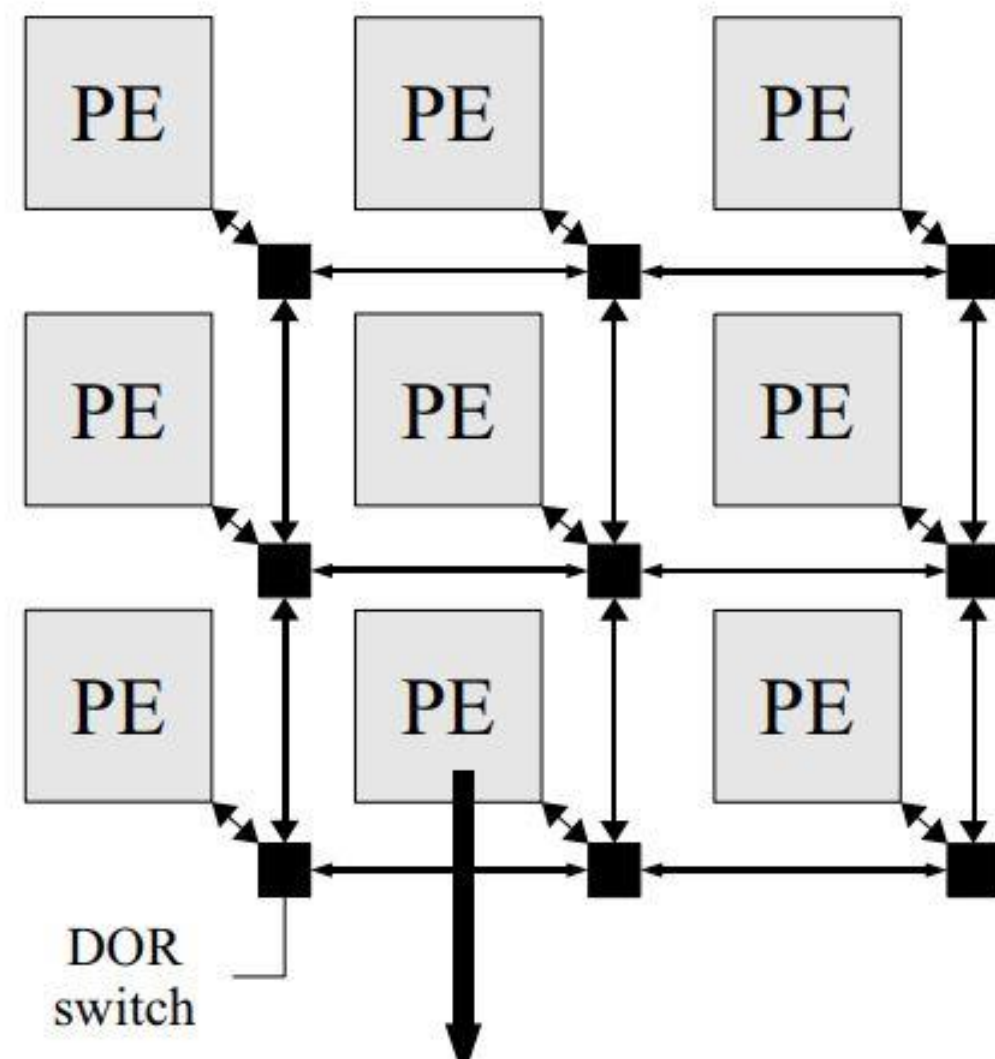


Grouping




Previous Work

T. Nechma and M. Zwolinski, "Parallel sparse matrix solution for circuit simulation on fpgas," IEEE Transactions on Computers, vol. 64, pp. 1090-1103, April 2015.






N. Kapre and A. DeHon, "Parallelizing sparse matrix solve for spice circuit simulation using fpgas," in 2009 International Conference on Field Programmable Technology, pp. 190-198, Dec 2009.




Results



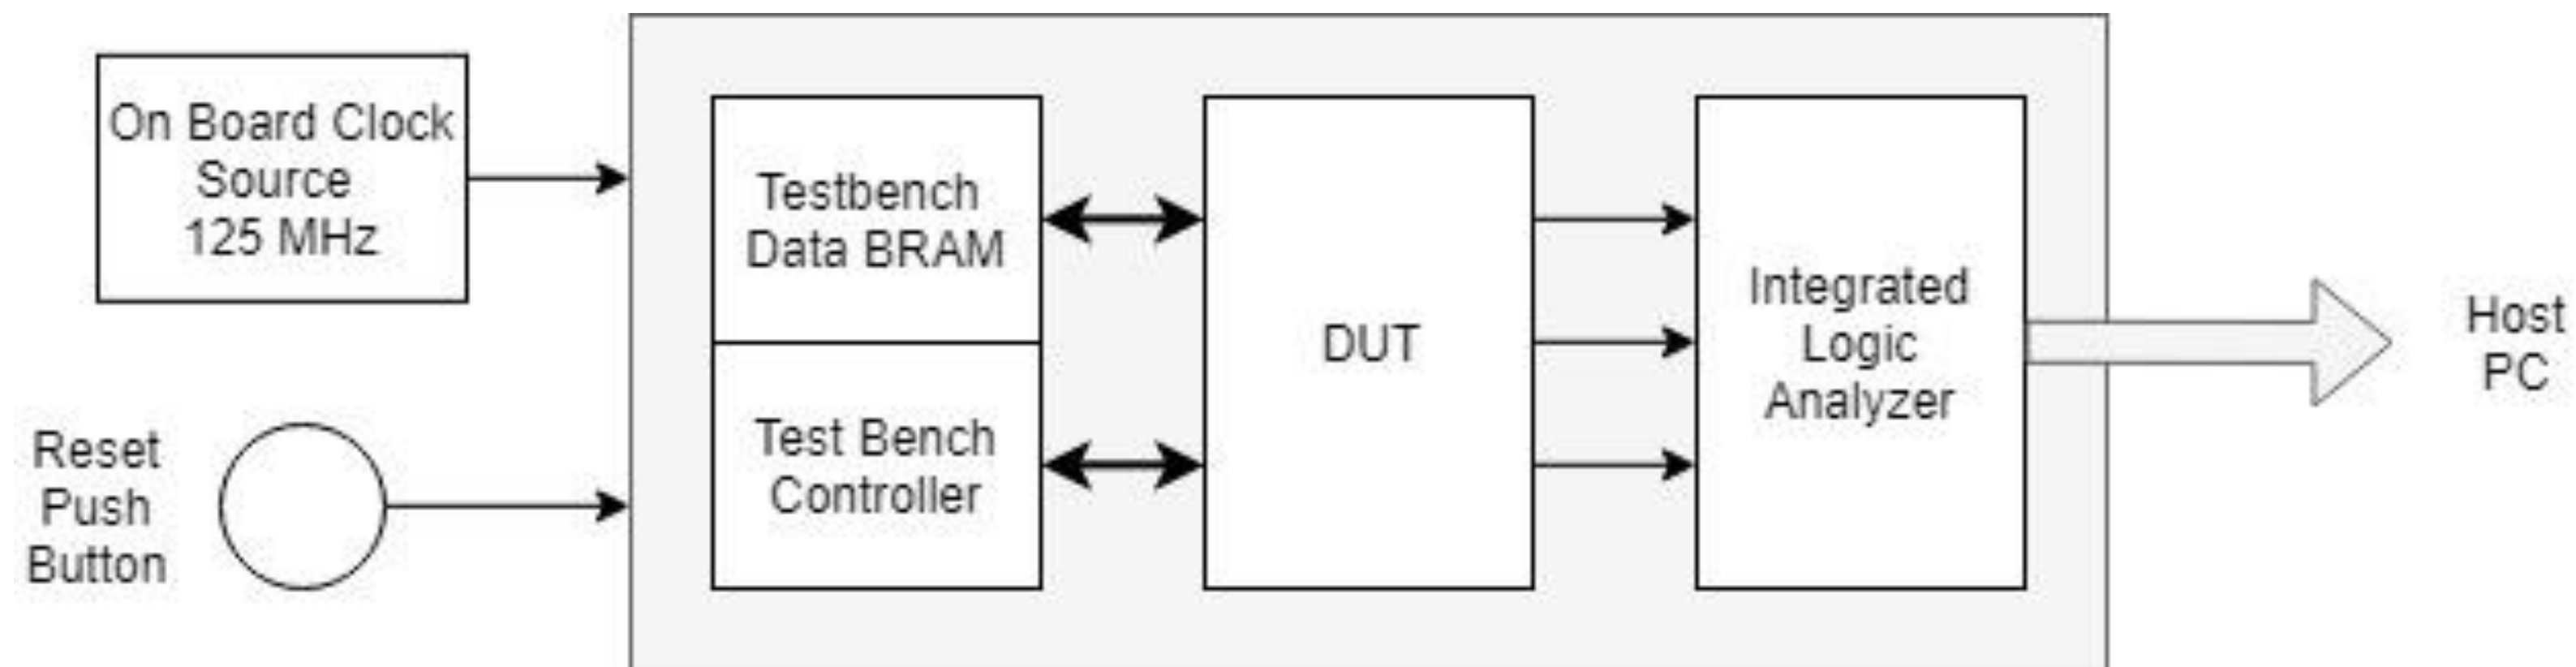
Performance Variation with the Number of Units



Performance Variation with the Latency of Units



Hardware Testing



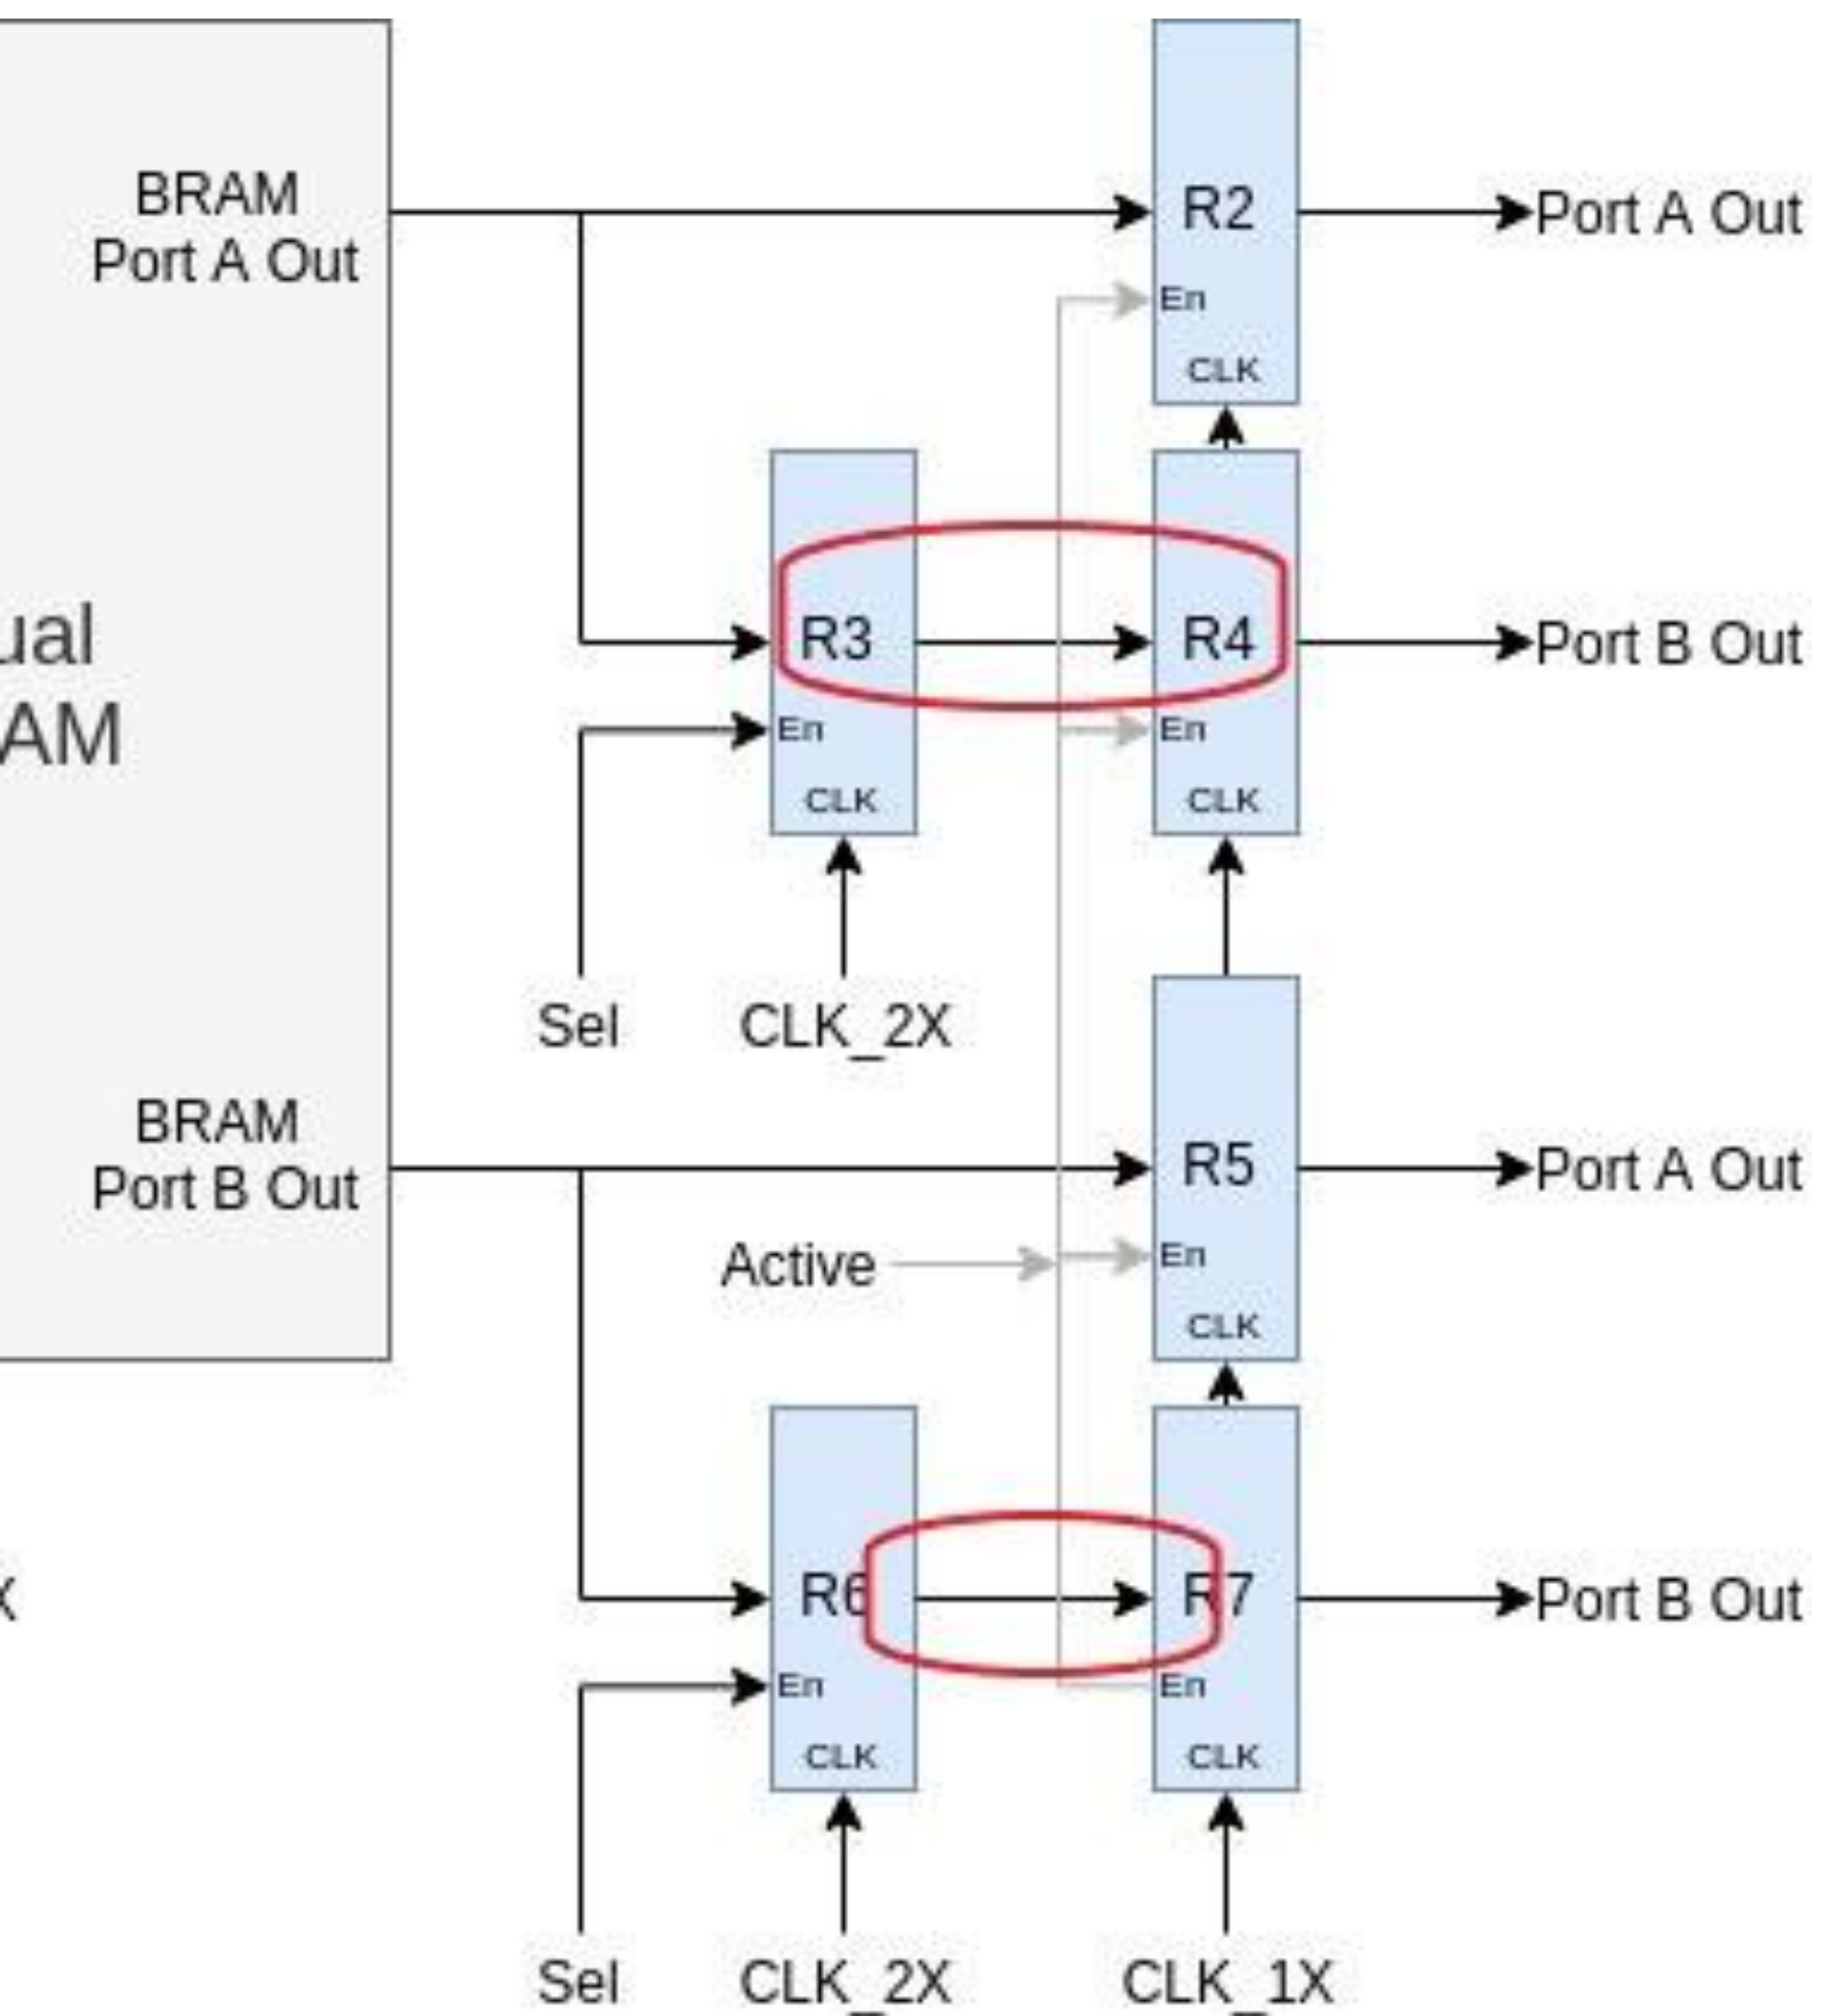
ILA ports are connected to the inputs and outputs of the BRAM ports

Hardware Configuration Parameter	Value
Number of PEs	4
Number of BRAMs	4
Number of Ports/ BRAM	4
Read Latency of BRAM	2
Latency of MAC Units	20
Latency of Diver Units	31

ILA Status: Idle												
Name	Value	3	4	5	6	7	8	9	10	11		
LUDH/hila/design_HILA_I/probe0_1	1											
> LUDH/hila/design.../probe1_1[42:0]	0007fc00000				0003ecdffef				0007fc00000			
> LUDH/hila/design.../probe2_1[42:0]	0007fc00000				0003ecdffef			0003ecd77ef	0007fc00000			
> LUDH/hila/design.../probe3_1[42:0]	0007fc00000	00040800000			0003ecdffef				0007fc00000			
> LUDH/hila/design.../probe4_1[42:0]	4024000bff5		0003ecdffef		4003ecdffef	0023ecdffef	0007fc00000	4024000bff5	0007fc00000	40...		
> LUDH/hila/design.../probe5_1[42:0]	0007fc00000				0003e4fbfbc				0007fc00000			
> LUDH/hila/design.../probe6_1[42:0]	0007fc00000				0003e4fbfbc			0003e451090	0007fc00000			
> LUDH/hila/design.../probe7_1[42:0]	4017fc00000		0003e4fbfbc		0023e4fbfbc	0003e4fbfbc	000c0800000	4017fc00000	0007fc00000			
> LUDH/hila/design.../probe8_1[42:0]	4037fc00000		0003e4fbfbc		4003e4fbfbc	0003e4fbfbc	0007fc00000	4037fc00000	0007fc00000			
> LUDH/hila/design.../probe9_1[42:0]	00040a00000				00040a00000							
> LUDH/hila/design.../robe10_1[42:0]	00040a00000				00040a00000							
> LUDH/hila/design.../robe11_1[42:0]	00040a00000				00040a00000							
> LUDH/hila/design.../robe12_1[42:0]	00040a00000				00040a00000	00040a00000	000c0801b10	00040a00000	40...			
> LUDH/hila/design.../robe13_1[42:0]	000be402d0b				000bf403600				000be402d0b			
> LUDH/hila/design.../robe14_1[42:0]	000be402d0b				000bf403600			000be402401	000be402d0b			
> LUDH/hila/design.../robe15_1[42:0]	000be402d0b	000bf403600	001bf403600	000bf403600	002c0a00000	000bf403600	00040c00000	000be402d0b	000be402d0b			
> LUDH/hila/design.../robe16_1[42:0]	000be402d0b	000bf403600	002bf403600	000bf403600	40040c00000	000bf403600		000be402d0b				
LUDH/hila/design_HILA_I/probe17_1	0											

Updated at: 2019-Jun-19 17:26:20

False Hold Violation




CLK_1X

CLK_2X

Sel

R3_Out[]



The path between R3 and R4 is not triggered at the rising edge of the CLK_1X hence the hold violation at this edge must be ignored



Future Work

- Synthesize BRAM
- Integrate with the on-chip ARM core
- New scheduling strategies

Thank you



Future Work

- Synthesize BRAM
- Port the tool on for on-chip ARM core
- New scheduling strategies

Parallelizing the Sparse LU Decomposition